



# AI 330: Machine Learning

## Fall 2023

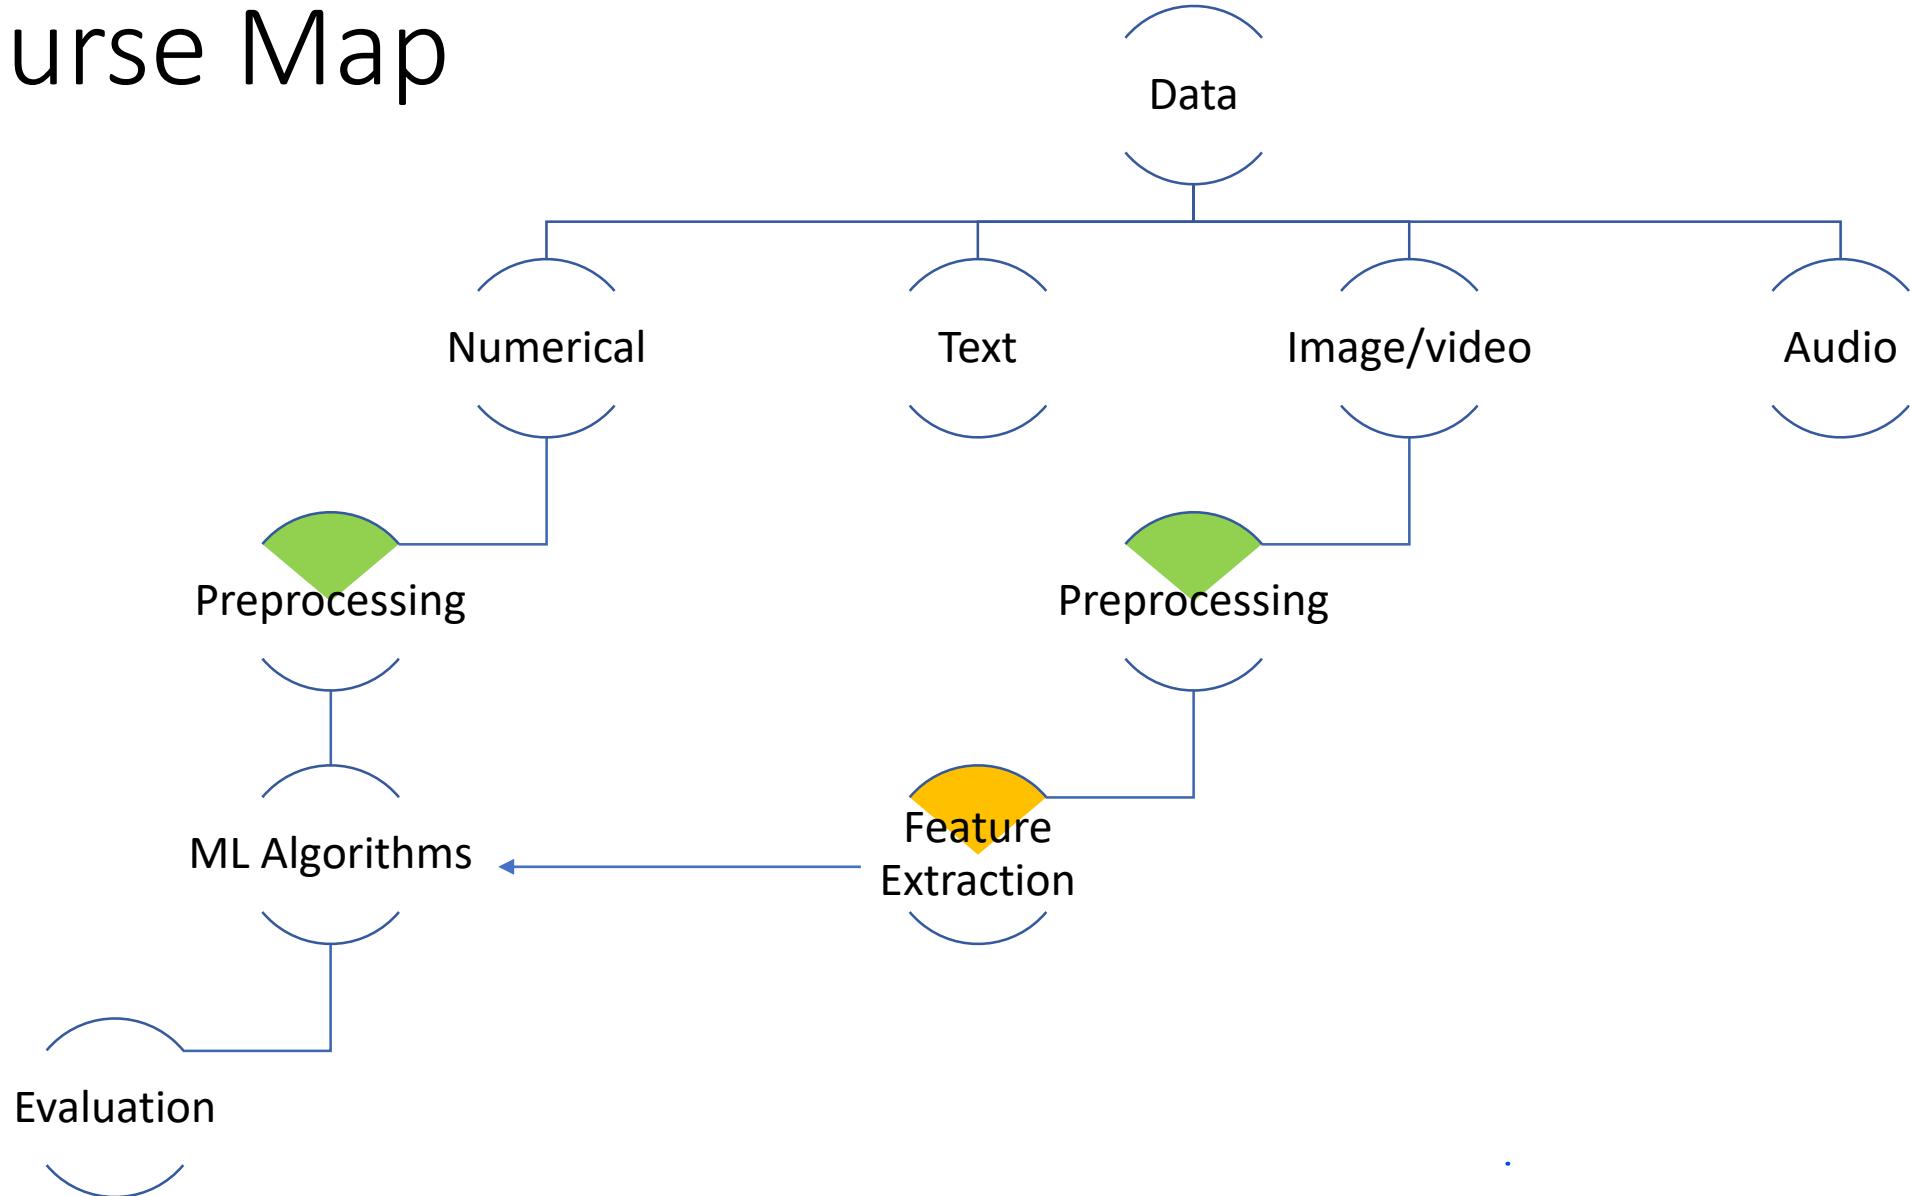
**Dr. Wessam EL-Behaidy**

Associate Professor, Computer Science Department,  
Faculty of Computers and Artificial Intelligence,  
Helwan University.

**Dr. Ensaif Hussein**

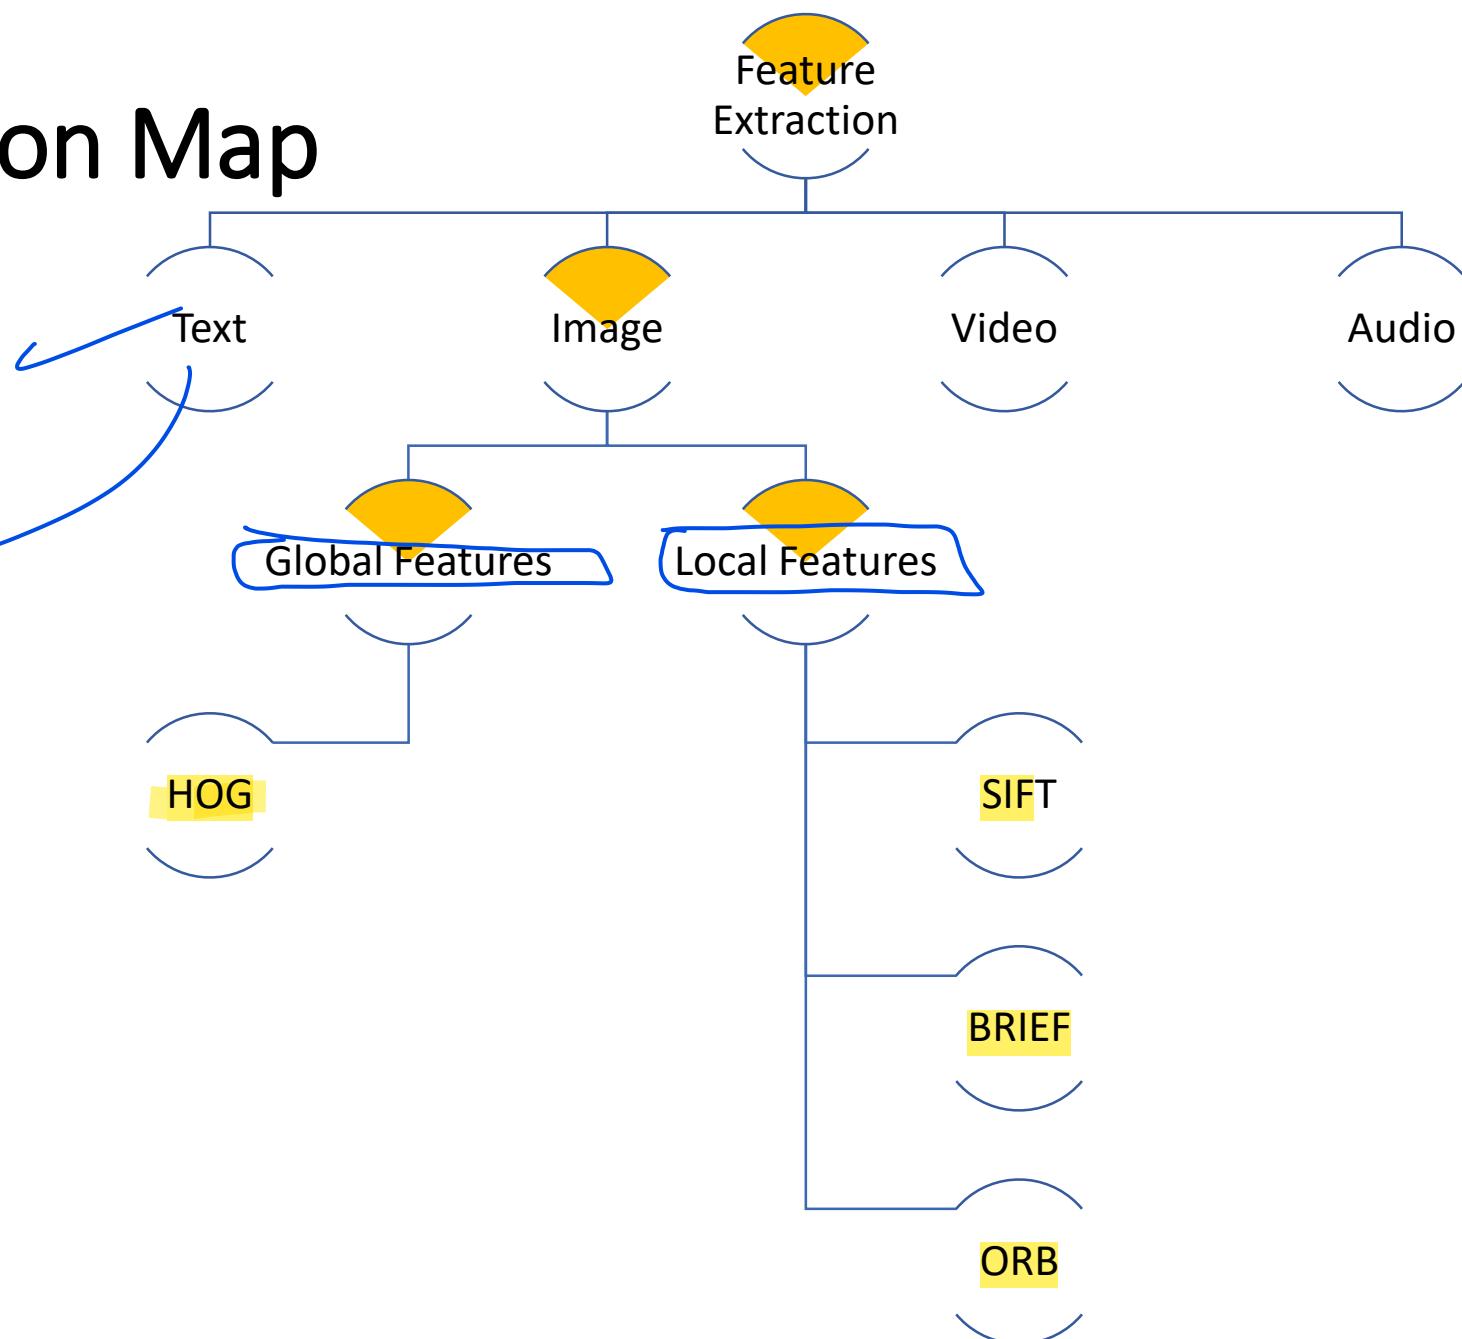
Associate Professor, Computer Science Department,  
Faculty of Computers and Artificial Intelligence,  
Helwan University.

# Course Map



# Feature Extraction Map

Feature Extraction



# Lecture 3

---

# Image Feature Extraction

# Feature Extraction

Transform the image into numerical data

- **Feature extraction** refers to the process of transforming raw data into numerical values (features).
- An image **feature descriptor** is a simplified representation of the image that contains only the most important information about the image, i.e., data unique to this specific object.
- Generally, features extracted from an image are of much lower dimension than the original image.
- A good feature descriptor is used to:
  - distinguish objects from one another. Ex: distinguish between cars and trucks.
  - match these features to a new image of the same object. Ex: recognize that this is the same building in two different images.

# Feature Extraction

Feature extraction can be accomplished manually or automatically:

- **Manual feature extraction (Hand-crafted features)** requires identifying and describing the features that are relevant for a given problem and implementing a way to extract those features. Over decades of research, engineers and scientists have developed feature extraction methods for images, signals, and text.

Manual feature extraction

Standard ML

- **Automated feature extraction** uses specialized algorithms or deep networks to extract features automatically from signals or images without the need for human intervention. This technique can be very useful when you want to move quickly from raw data to developing machine learning algorithms.

one step

SDL

Feature extraction remains the first challenge that requires significant expertise before one can build effective predictive models.

# Image Feature Descriptor

- **Global features** describe the image as a whole to generalize the entire object.
- Whereas the **local features** describe the image patches (key points in the image) of an object.
- Generally,
  - for low-level applications such as object detection, global features are used.
  - for higher-level applications such as object recognition, local features are used.

الرسومات اليدوية المكتوبة باللغة العربية:

- التفاصيل العامة ت 描 ع الصورة ككل ل-generalization.
- التفاصيل المحلية ت 描 ع الأجزاء الصغيرة من الصورة (الpatches) كpoints key points.
- عادةً ما يتم استخدام التفاصيل العامة في تطبيقات المستوى المنخفض مثل التعرف على الأشياء.
- عادةً ما يتم استخدام التفاصيل المحلية في تطبيقات المستوى العالي مثل التعرف على الأشياء.

# Global & local features

- There are number of features for images, the most popular ones are:
  - Global features:
    - HOG – Histogram of Oriented Gradients
  - Local features:
    - SIFT – Scale Invariant Feature Transform
    - BRIEF – Binary Robust Independent Elementary Features
    - ORB – Oriented FAST Rotated BRIEF

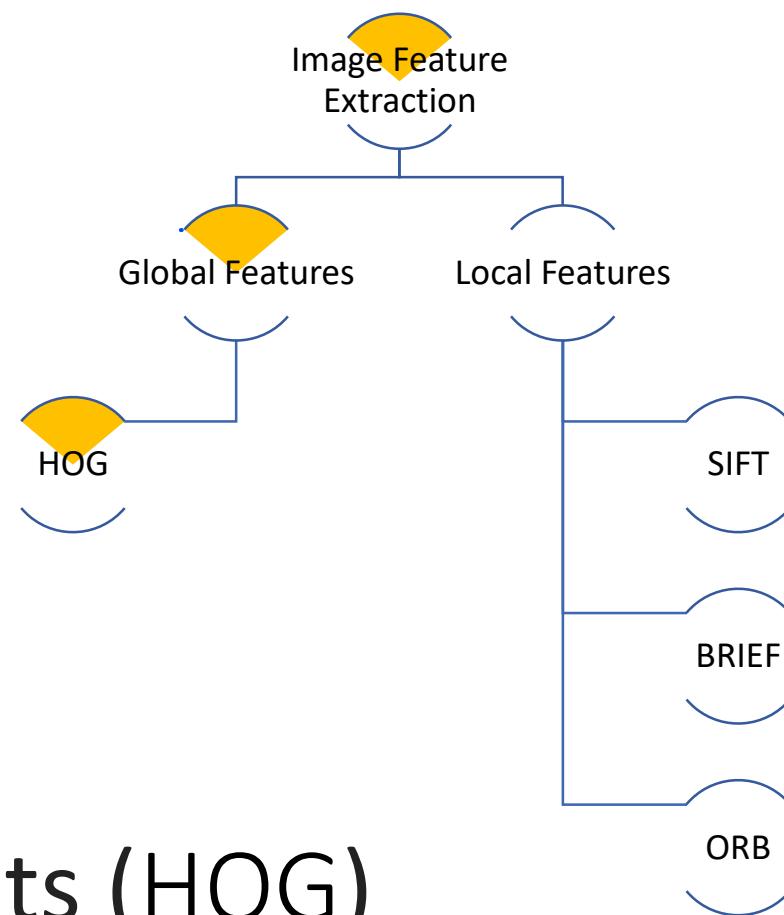
اٹو سی جو ہر جیسے  
جیسا کو ایک ہدایت کرو دیتے ہیں

## Global Feature:

### Histogram of Oriented Gradients (HOG)

#### Slides References:

- <https://www.analyticsvidhya.com/blog/2019/09/feature-engineering-images-introduction-hog-feature-descriptor/#:~:text=HOG%2C%20or%20Histogram%20of%20Oriented,vision%20tasks%20for%20object%20detection.>
- <https://learnopencv.com/histogram-of-oriented-gradients/>



# Histogram of Oriented Gradients

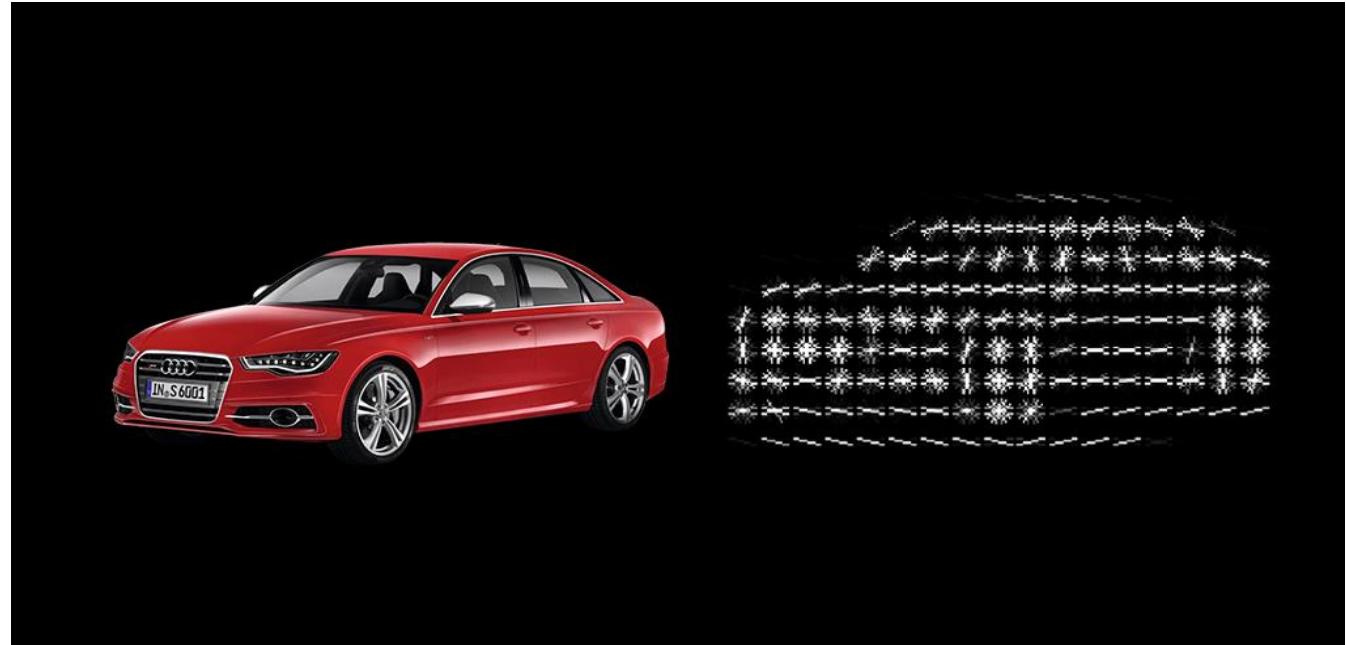
- HOG, or Histogram of Oriented Gradients, is a feature descriptor that is often used to extract **global features** from image data.
- It is widely used in **computer vision** tasks for **object detection**.

intelligent / jhwo's ↪

# Histogram of Oriented Gradients

- The HOG descriptor:
  - focuses on the structure or the shape of an object. HOG is able to identify the edges and provide the edge direction as well. This is done by extracting the **gradient and orientation** (or you can say magnitude and direction) of the edges

*Gradient  $\rightarrow$  Magnitude  
Gradient  $\rightarrow$  direction*



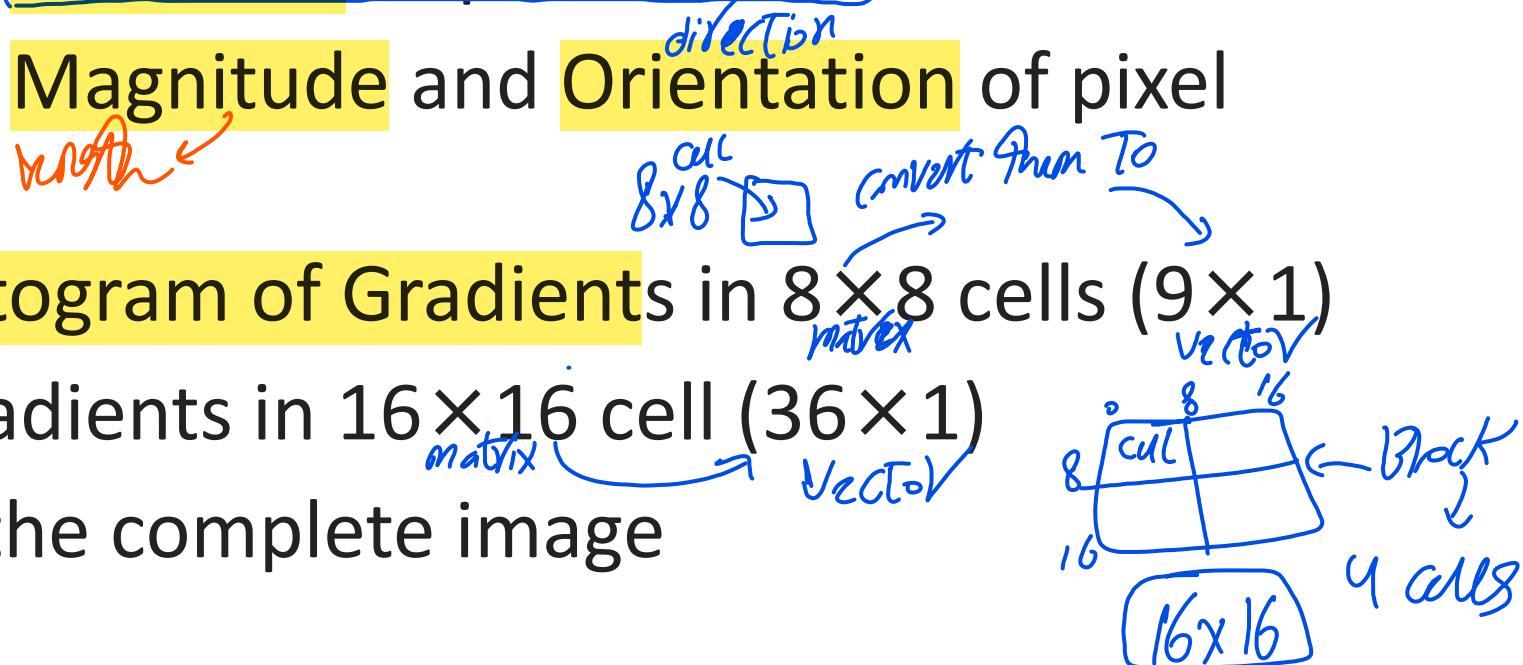
A formal definition:

The HOG feature descriptor counts the occurrences of gradient orientation in localized portions of an image.

# HOG Steps

الخطوات الستة في HOG  
 Resize always

- Step 1: Preprocessing the image (Crop and resize).
- Step 2: Calculate the **Gradients** of pixel values.
- Step 3: Calculate the **Magnitude** and **Orientation** of pixel values.
- Step 4: Calculate **Histogram of Gradients** in  $8 \times 8$  cells ( $9 \times 1$  matrix)
- Step 5: Normalize gradients in  $16 \times 16$  cell ( $36 \times 1$  vector)
- Step 6: Features for the complete image



# Step 1: Preprocessing the image

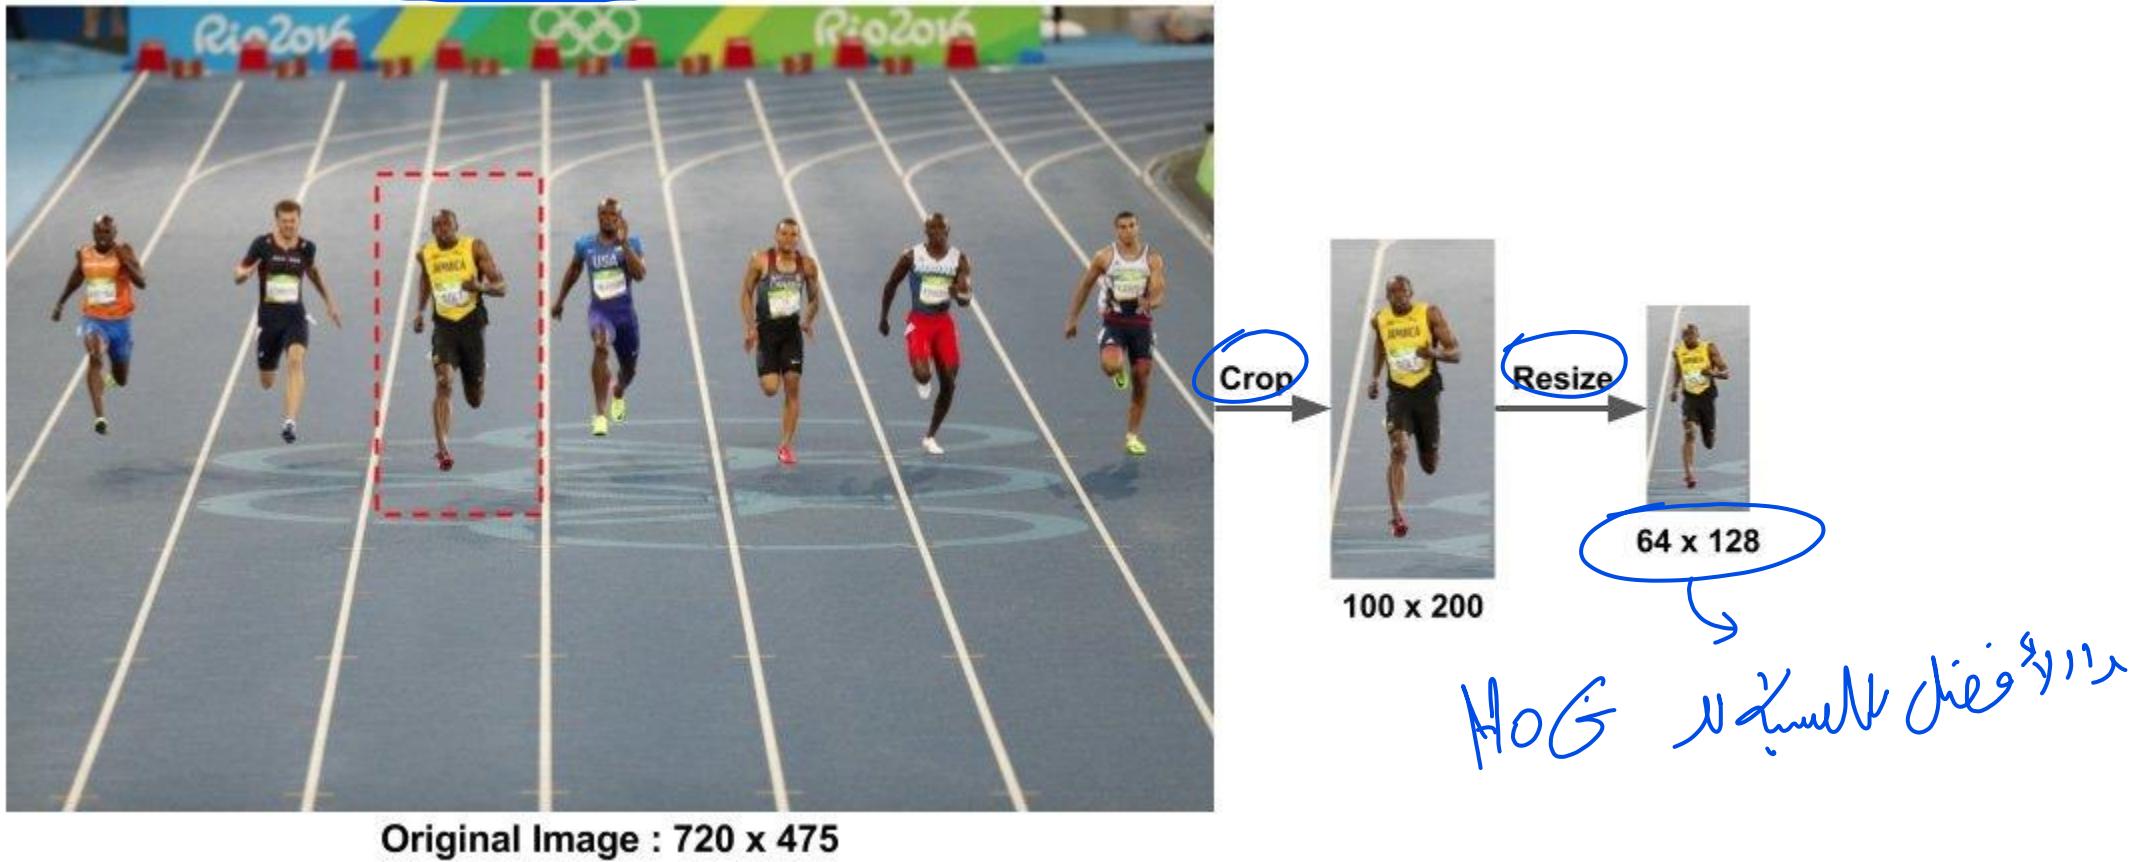
- HOG feature descriptor is calculated on a  $64 \times 128$  patch of an image.
- Of course, an image may be of any size.
- Typically patches at multiple scales are analyzed at many image locations. The only constraint is that the patches need to have an aspect ratio of  $1:2$ 
  - For example, they can be  $100 \times 200$ ,  $128 \times 256$ , or  $1000 \times 2000$  but not  $101 \times 205$ .

$1:2$      $1:2$      $1:2$

~~$101:205$~~      $105:210$  ✓

# Step 1:

- Example: We used an image of size  $720 \times 475$ . We selected a patch of size  $100 \times 200$  for calculating our HOG feature descriptor. This patch is cropped out of an image and resized to  $64 \times 128$



## Step 2: Calculate the Gradients

جیز اور خود کیا

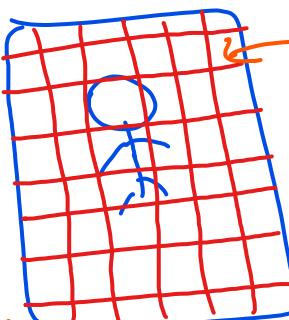
- **Gradients** are the small change in the x and y directions.
- We calculate the gradients on each pixel in the image, by filtering it with a kernel, like **Sobel** with size 1.
- To determine the gradient (or change) in the x-direction

نیشن  
جیز اور خود کیا

Ex: Change in X direction ( $G_x$ ) =  $-78 + 89 = 11$

- Similarly, to calculate the gradient in the y-direction

Ex: Change in Y direction ( $G_y$ ) =  $-76 + 68 = -8$

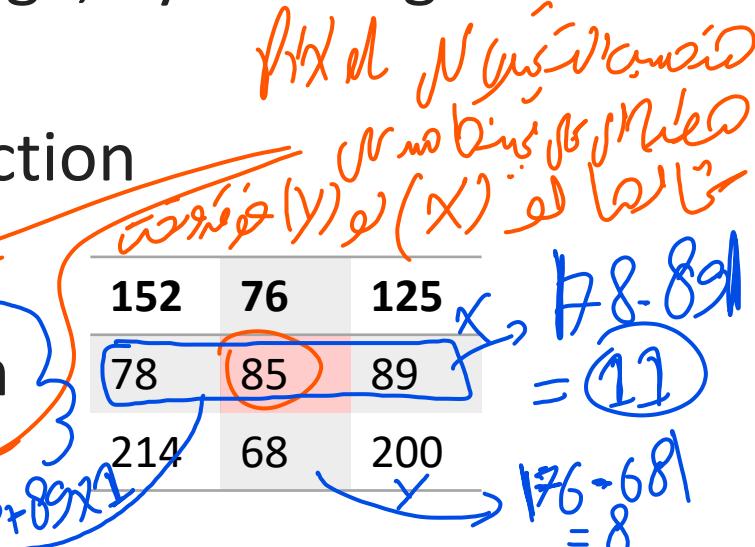


Gradient, fixd Nisian  
8x8 fixd Nisian  
Grad N  
absdilu pbs g fixd  
90W, SD all  
J 90W, SD all  
8x8  
Sobel kernel of size 1  
Yi x i fixd Nisian

-1	0	1
0		
1		

Sample of image pixel values

Gradient calculated  
Pixel magnitude  
 $\theta$  (in direction)  
 $m = \sqrt{(G_x)^2 + (G_y)^2}$   
 $\theta = \tan^{-1} \left( \frac{G_y}{G_x} \right)$



## Step 2:

- We have calculated the gradients in both x and y directions separately. So, this process will give us two new matrices:

- one storing gradients in the x-direction and
- the other storing gradients in the y direction.

## Step 3: Calculate the Magnitude and Orientation (direction)

- Using the gradients we calculated in the last step, we will now determine the **magnitude** and **direction (orientation)** for each pixel value using these formulas:

$$\text{Magnitude} = \sqrt{(G_x)^2 + (G_y)^2}$$

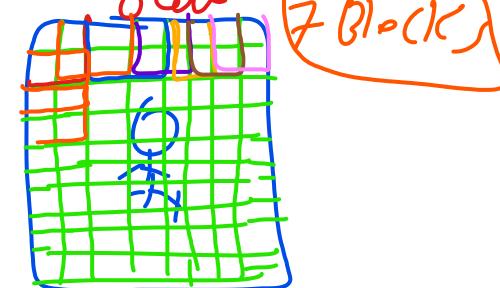
*16x8 cells, 16x8 block, New size 16x8, 16x8 gradient direction*

$$\Phi = \arctan(G_y / G_x)$$

- So, for the previous example, we had  $G_x$  and  $G_y$  as 11 and 8.

$$\text{Magnitude} = \sqrt{(11)^2 + (8)^2} = 13.6$$

$$\Phi = \arctan(8 / 11) = 36$$



Total directions = 160000 x 36 = 3750

Total Blocks = 16 x 7 = 112 Block

For color images, the gradients of the three channels are evaluated. The magnitude of gradient at a pixel is the maximum of the magnitude of gradients of the three channels, and the angle is the angle corresponding to the maximum gradient

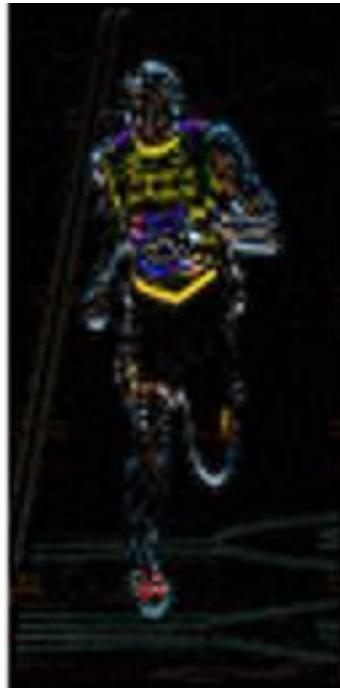
Block size = 3x3 histogram bin  
9 histogram bins in one cell  
11M cells in one block

## Step 3:

Absolute value of x-gradient



Absolute value of y-gradient



Magnitude of gradient



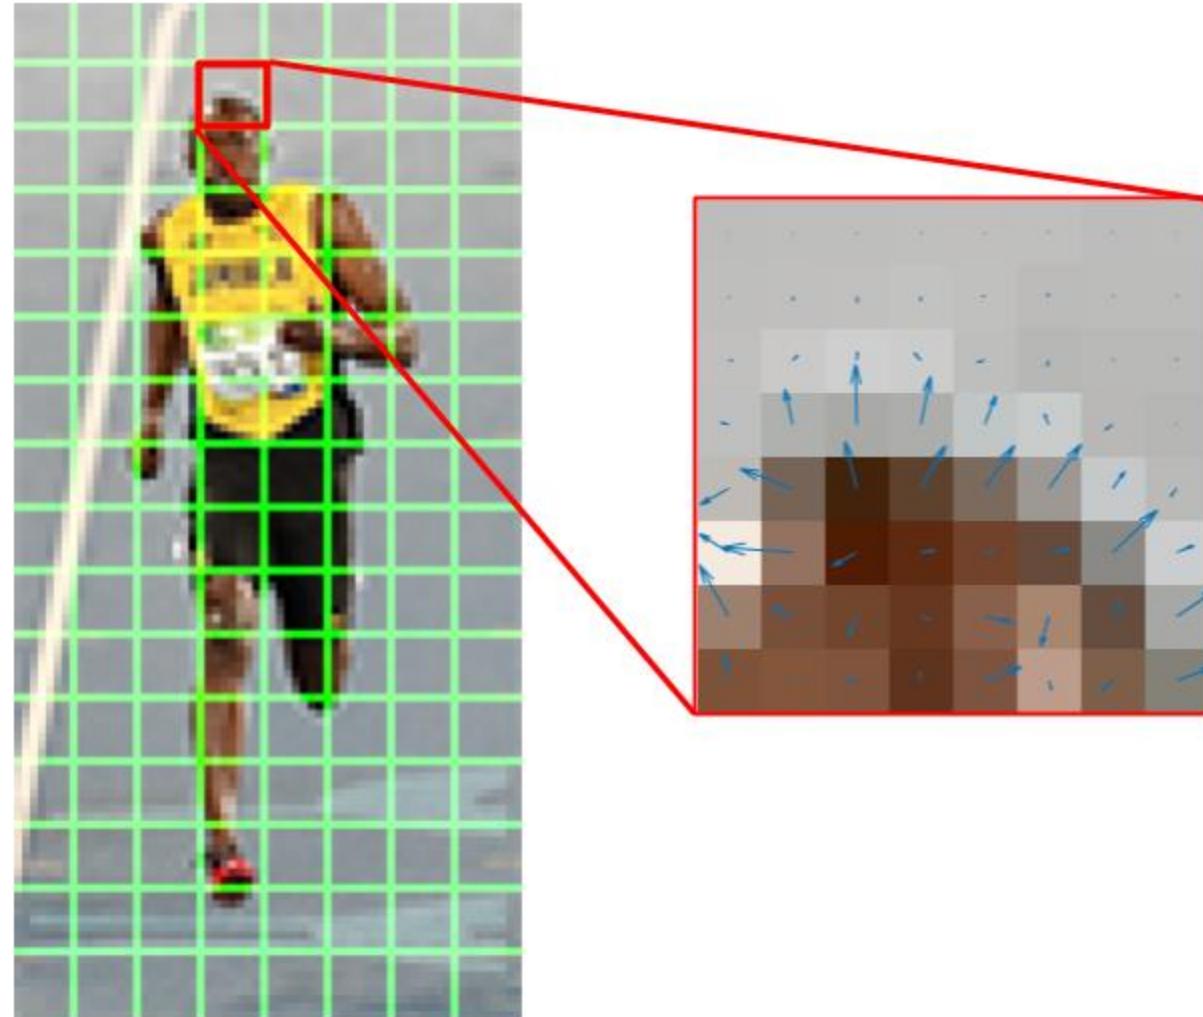
- Notice, that the x-gradient fires on vertical lines and the y-gradient fires on horizontal lines. The magnitude of gradient fires where ever there is a sharp change in intensity (as edges). None of them fire when the region is smooth.
- The gradient image removed a lot of non-essential information ( e.g. constant colored background ) but highlighted outlines

## Step 4: Calculate Histogram of Gradients in $8 \times 8$

cells ( $9 \times 1$ )  $\xrightarrow{(9 \times 1) \text{ bins for each Pixel}}$

$8 \times 8 \rightarrow 256 \text{ bins}$

- In this step, the image is divided into  $8 \times 8$  cells and a histogram of gradients is calculated for each  $8 \times 8$  cell.
- We can certainly change this value here from  $8 \times 8$  to  $16 \times 16$  or  $32 \times 32$ .
- In the figure, the **direction of arrows** points to the direction of change in intensity, and the **magnitude** shows how big the difference is.



(8x8) direction by magnitude 11 گوئی ڈبلس 11 یوں ہے 8x8

8x8

## Step 4:

- We have 9 bins in the histogram.
- We select:

- bin**: based on the direction
- vote** (the value that goes into the bin): based on the magnitude.

Gradient by change in direction  
Y → 11 11 Y 11

- Ex1:** the pixel is encircled in blue. angle =  $80^\circ$ , magnitude = 2. angle (direction) =  $80^\circ$ , magnitude = 2.

→ So it adds 2 to the 5th bin.

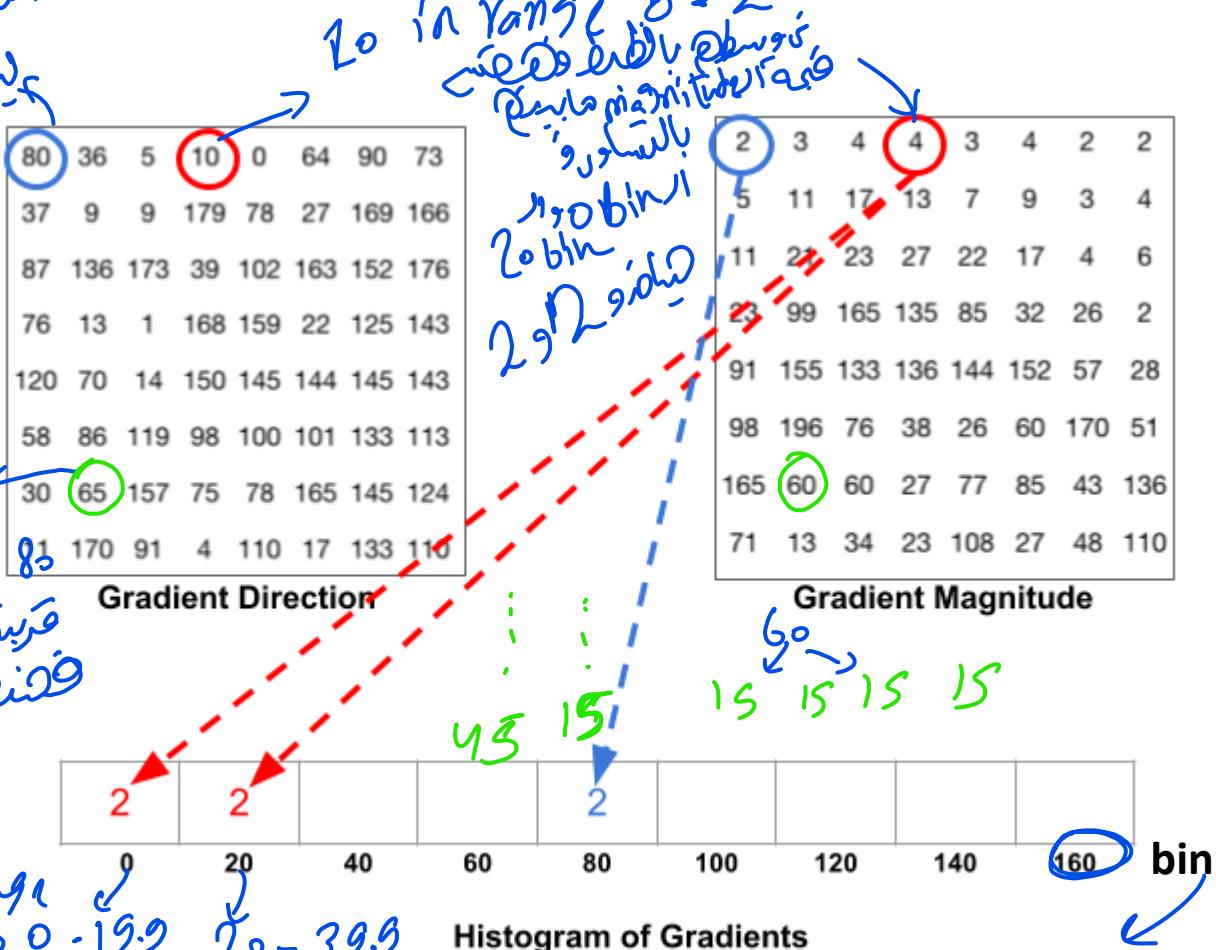
- Ex2:** The pixel is encircled in red. angle =  $10^\circ$ , magnitude = 4.

Since  $10^\circ$  is halfway between 0 and 20,

→ the vote by the pixel splits evenly into the two bins

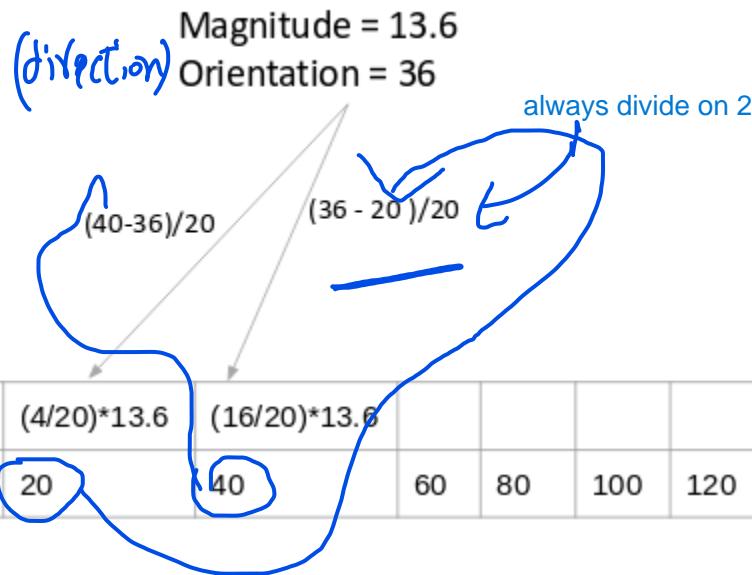
10 is in 1 bin 10 is in 2 bins  
for each cell

جس کا Use (9x1)  
کر دیں



9 bins x 20  
20 اسی Range 1 bin میں

# Step 4:



Ex3: Uneven split in direction (orientation)

80	36	5	10	0	64	90	73
37	9	9	179	78	27	169	166
87	136	173	39	102	163	152	176
76	13	1	168	159	22	125	143
120	70	14	150	145	144	145	143
58	86	119	98	100	101	133	113
30	65	157	75	78	165	145	124
11	170	91	4	110	17	133	110

**Gradient Direction**

$$(165-160)/20$$



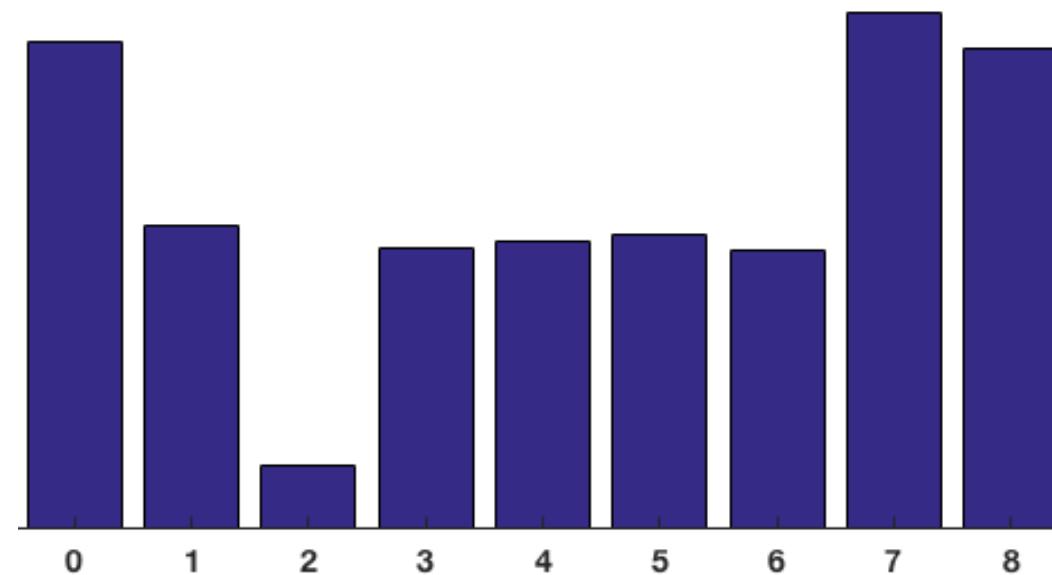
Histogram of Gradients

Ex4: If the angle is greater than 160 degrees

**Remember,** the higher contribution should be to the bin value which is closer to the orientation

## Step 4:

- The contributions of all the pixels in the  $8 \times 8$  cells are added up to create the 9-bin histogram.
- we will get a  $9 \times 1$  matrix for each cell, and its 9-bin histogram looks like this:



ow (36x1) gradients block's 4 cell for 16x16 cell 16x16

## Step 5: Normalize gradients in 16×16 cell (36×1)

16x16 is zero, now overall use 8x8

- In the previous step, we created a histogram based on the gradient of the image.
  - Gradients of an image are sensitive to overall lighting.
- Ideally, we want our descriptor to be independent of lighting variations.
  - We can reduce this lighting variation by “normalizing” the gradients over a bigger sized block of 16×16.

## Step 5:

- To perform a **16 x 16 block**, we will be combining four  $8 \times 8$  cells to create a  $16 \times 16$  block.
  - We have four  $9 \times 1$  matrices (  $9 \times 1$  matrix from  $8 \times 8$  cell) → 1 cell
  - A  $16 \times 16$  block can be concatenated to form a  $36 \times 1$  matrix
- To normalize this matrix, we will divide each of these values by the square root of the sum of squares of the values.

Example,

for a given vector V:  $V = [a_1, a_2, a_3, \dots, a_{36}]$

We calculate the root of the sum of squares:

$$k = \sqrt{(a_1)^2 + (a_2)^2 + (a_3)^2 + \dots + (a_{36})^2}$$

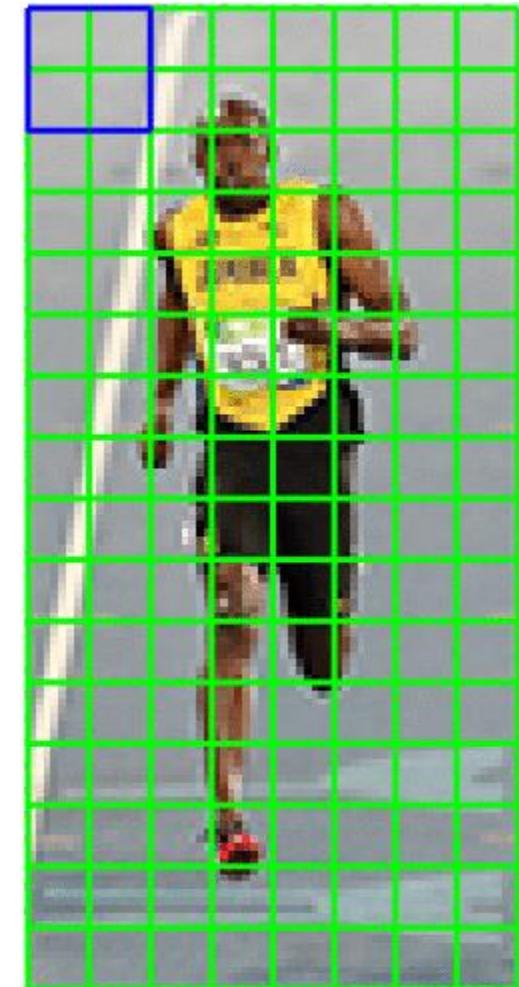
And divide all the values in the vector V with this value k:

Normalised Vector =  $\left( \frac{a_1}{k}, \frac{a_2}{k}, \frac{a_3}{k}, \dots, \frac{a_{36}}{k} \right)$

The resultant would be a **normalized vector of size  $36 \times 1$**

## Step 5:

- The window is then moved by 8 pixels (see animation ) and a normalized  $36 \times 1$  vector is calculated over this window and the process is repeated.

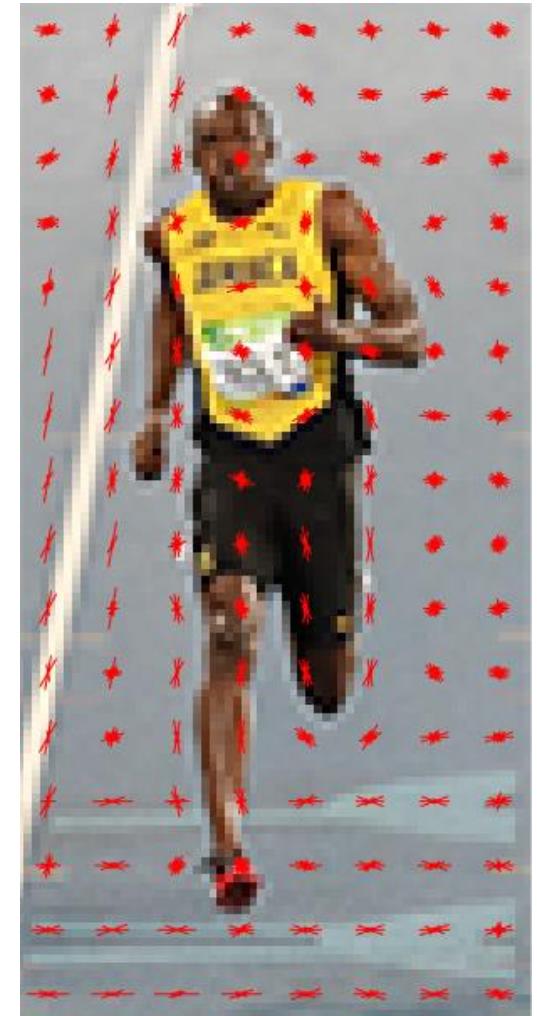


# Step 6: Features for the complete image

- To calculate the final feature vector for the entire image patch, the  $36 \times 1$  vectors are concatenated into one giant vector.
- Let us calculate  
1. How many positions of the  $16 \times 16$  blocks do we get for a single  $64 \times 128$  image?  
There are 7 horizontal and 15 vertical positions  $\rightarrow 7 \times 15 = 105$  positions.  
2. Each  $16 \times 16$  block is represented by a  $36 \times 1$  vector.  
 $\rightarrow$  The feature vector =  $36 \times 105 = 3780$  dimensional-vector.
- Hence, the total features for the image would be  $3780$  features.

# Visualizing Histogram of Oriented Gradients

- The HOG descriptor of an image patch is usually visualized by plotting the  $9 \times 1$  normalized histograms in the  $8 \times 8$  cells.
- See the image on the side. You will notice that the dominant direction of the histogram captures the shape of the person, especially around the torso and legs.

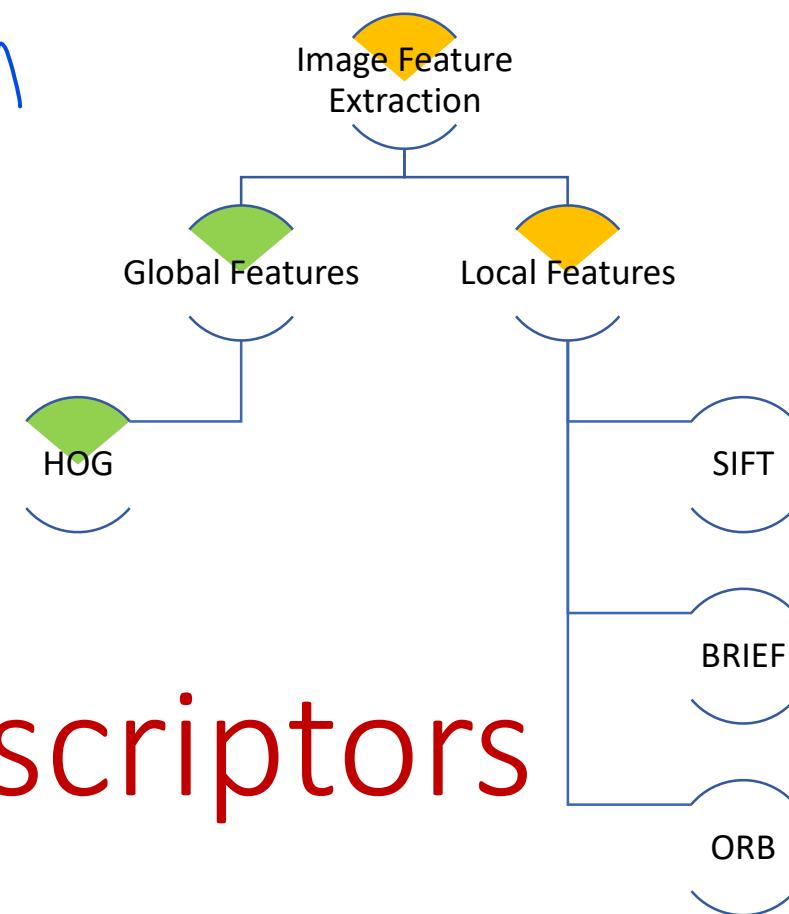


Numerical Vector into ML Model  
Classifier  
Feature abstraction

Significant object features

## Local features:

# Keypoints & Descriptors



### Slides References:

<https://www.analyticsvidhya.com/blog/2019/10/detailed-guide-powerful-sift-technique-image-matching-python/>

### University of Bonn , Prof. Cyrill Stachniss:

<https://www.youtube.com/watch?v=nGya59Je4Bs&t=607s>

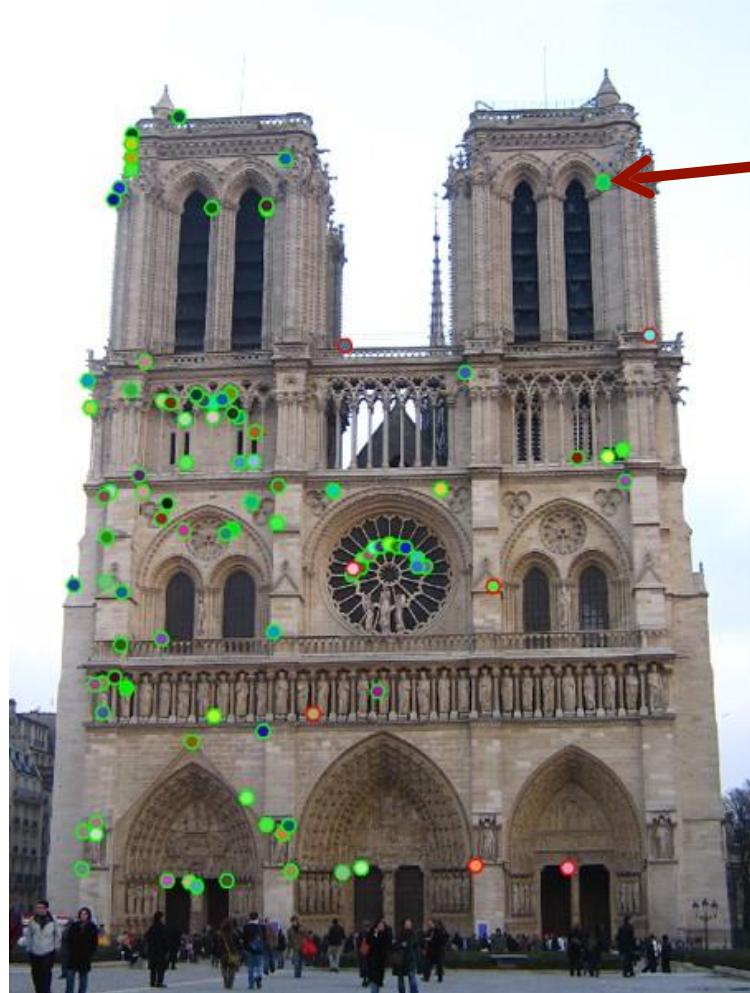
<https://www.youtube.com/watch?v=CMolhcwtGAU>

# Motivation

- Finding locally distinct points  
→ keypoints



# Keypoints and Descriptors



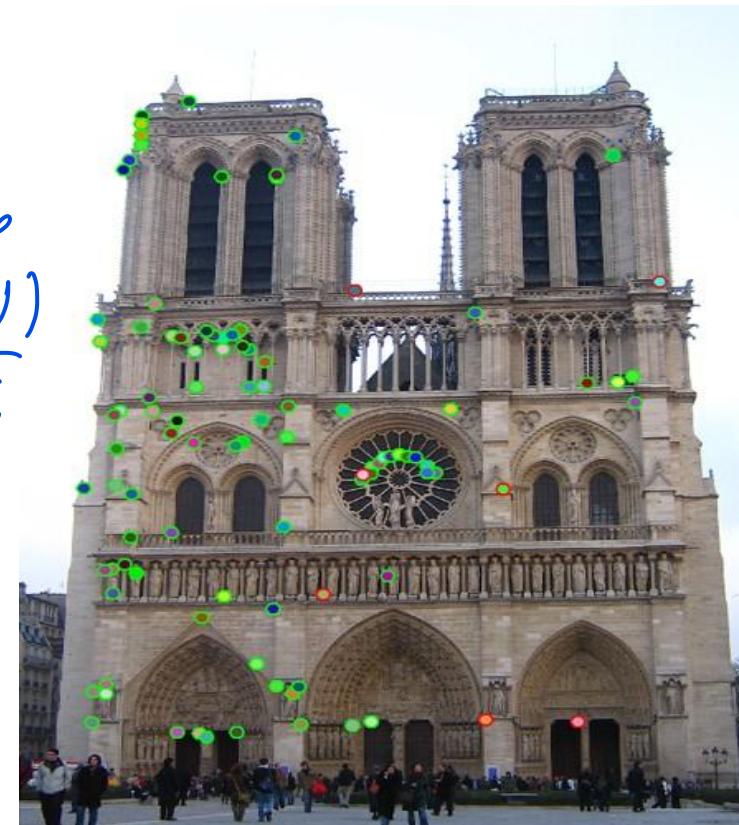
**keypoint**  
**descriptor** at  
the keypoint

$$f = \begin{bmatrix} 0.02 \\ 0.04 \\ 0.1 \\ 0.03 \\ 0 \\ \dots \end{bmatrix}$$

# Keypoints and Descriptors

- § **Keypoint** is a (locally) distinct location in an image
- § The feature **descriptor** summarizes the local structure around the keypoint

Keypoint  
+ *nichts*  
Descriptor  
= visual features



# Keypoints:

- § Corners are often highly distinct points
- § Corners are invariant to translation, rotation, and illumination.
- § Corner = two edges in roughly orthogonal directions
- § Edge = a sudden brightness change

وكلمة KeyPoint

مُنْبَهٌ لِّعَرْقِيَّةِ

لِّعَرْقِيَّةِ

لِّعَرْقِيَّةِ

لِّعَرْقِيَّةِ

لِّعَرْقِيَّةِ

لِّعَرْقِيَّةِ

لِّعَرْقِيَّةِ

# Keypoints:

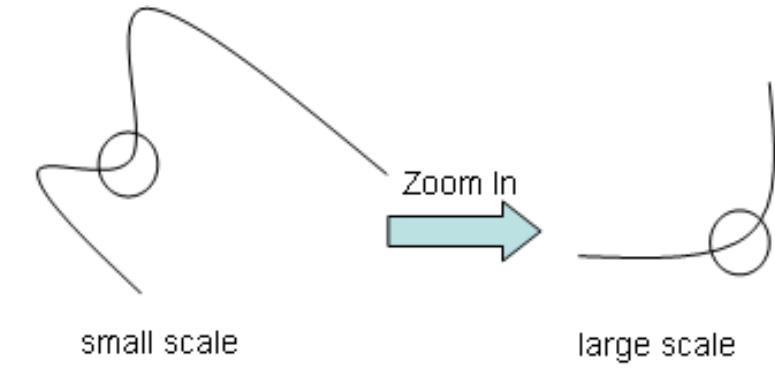
- § Two groups of approaches for finding locally distinct points:

## 1. Corners via structure matrix

§ Harris, Shi-Tomasi, Förstner

## 2. Difference of Gaussians (DoG)

§ Finds corners and blobs over scales



Scaling Problem in Harris

- § These approaches are key ingredients of most hand-designed features

Scale

Als den COrner nicht einzuordnen ist, kann man  
die Größe des COrners abstimmen.  
Die Größe des COrners ist abhängig von der  
Strukturmatrix.

# Local features: Keypoints & Descriptors

Difference of Gaussians (DoG) Keypoint Detection

# Difference of Gaussians Keypoints

z463 flev. DoG 11

- § A variant of corner detection
- § Provides responses at corners, edges, and blobs
- § **Blob** = mainly constant region but different to its surroundings



# Keypoints: DoG Over Scale-Space Pyramid

**Procedure** Over different image pyramid levels

**Step 1:** Gaussian smoothing (blurring)

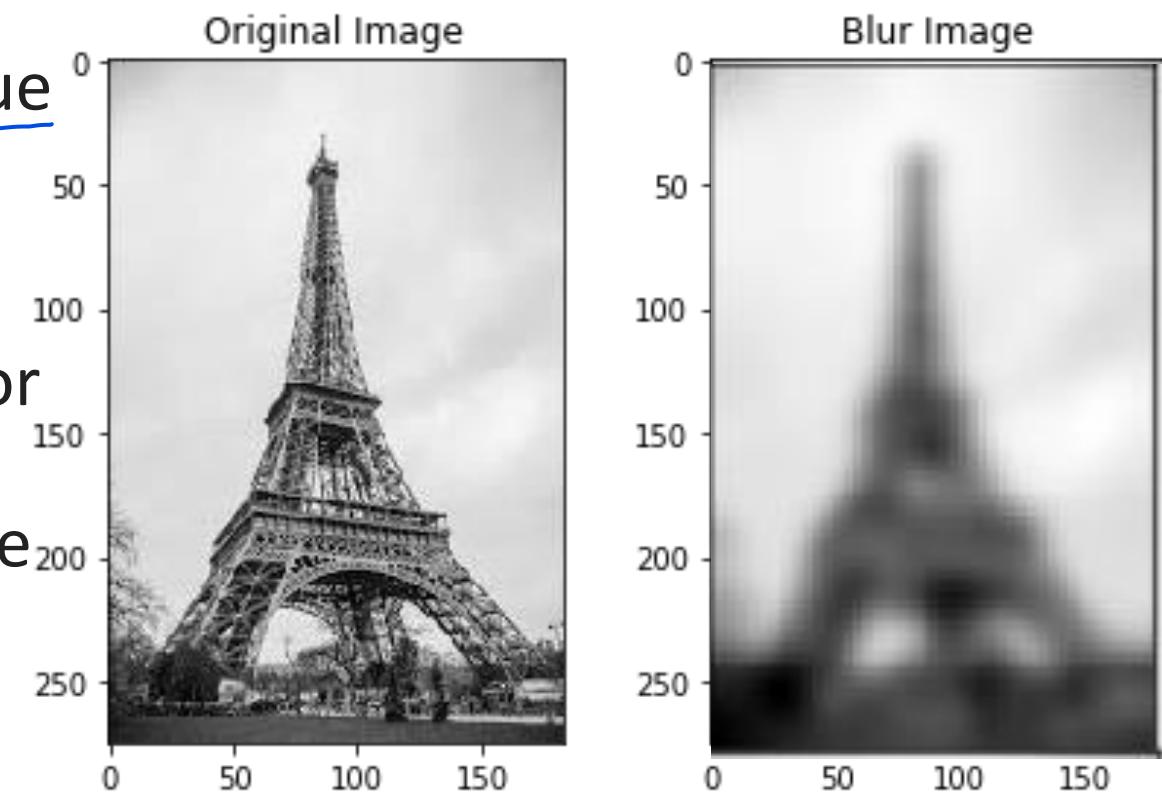
*(all pyramid Smoothing steps)*

**Step 2:** Difference-of-Gaussians: find extrema (over smoothing scales)

**Step 3:** maxima suppression at edges (i.e., **Keypoint Localization**)

# Step 1: Gaussian smoothing (Blurring)

- For every pixel in an image, the Gaussian smoothing calculates a value based on its neighboring pixels.
- As you can see, the texture and minor details are removed from the image, and only the relevant information like the **shape** and **edges** remain.



<https://www.analyticsvidhya.com/blog/2019/10/detailed-guide-powerful-sift-technique-image-matching-python/>

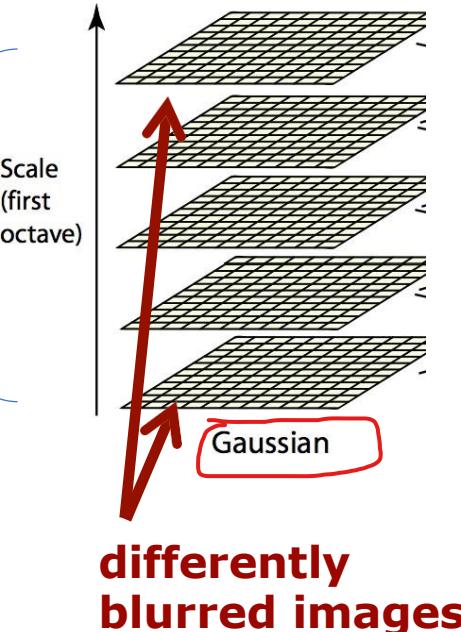
اے لیلیک وائے جیس جھوٹیں جیں  
(Shaff - 2d9E) ۱۱۰، ۹۰۰، ۱۱۰

# Illustration

- Gaussian smoothing reduces noise and highlights the important features of the image.
- Blur images with different  $\sigma$

different  
bands

**octave**



Blurring نیویوسی  
بلور  
Scaling نیویوسی  
موجاکتیوی

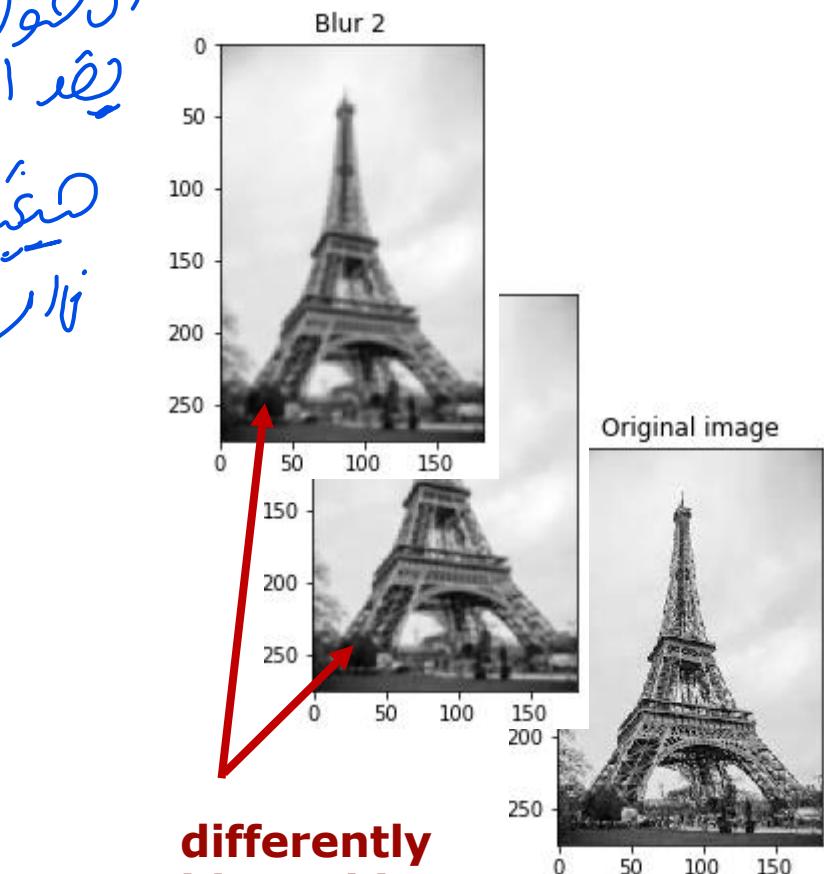


Image courtesy: Lowe38

# Scale Space

Different sizes of gradients  
Structure matrix values will increase.

Now, we need to ensure that these features must not be scale-dependent

→ searching for these features on multiple scales.

“**Scale space**” is a collection of images having different scales, generated from a single image.

# Illustration

with different scales (scales) get images size 2nd octave is  
also by blurring (smoothing) always next octave is

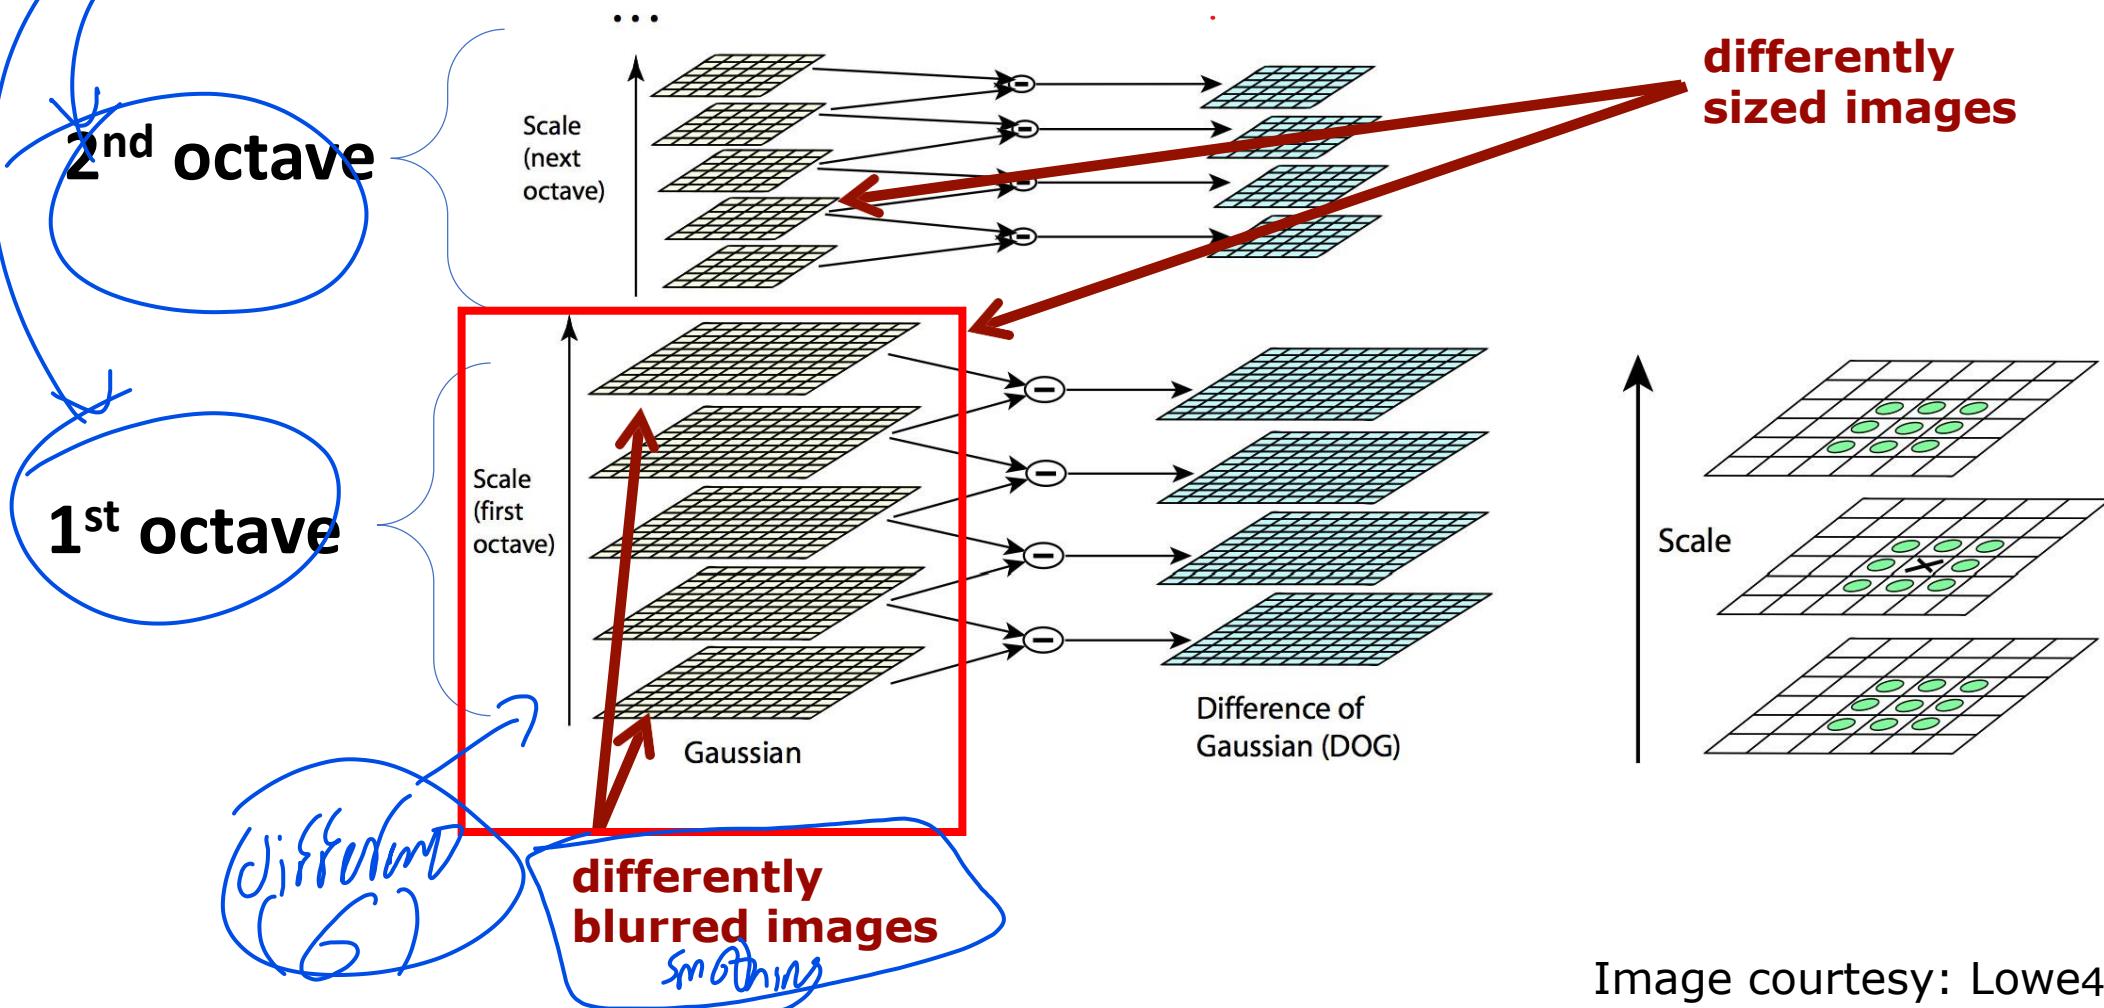
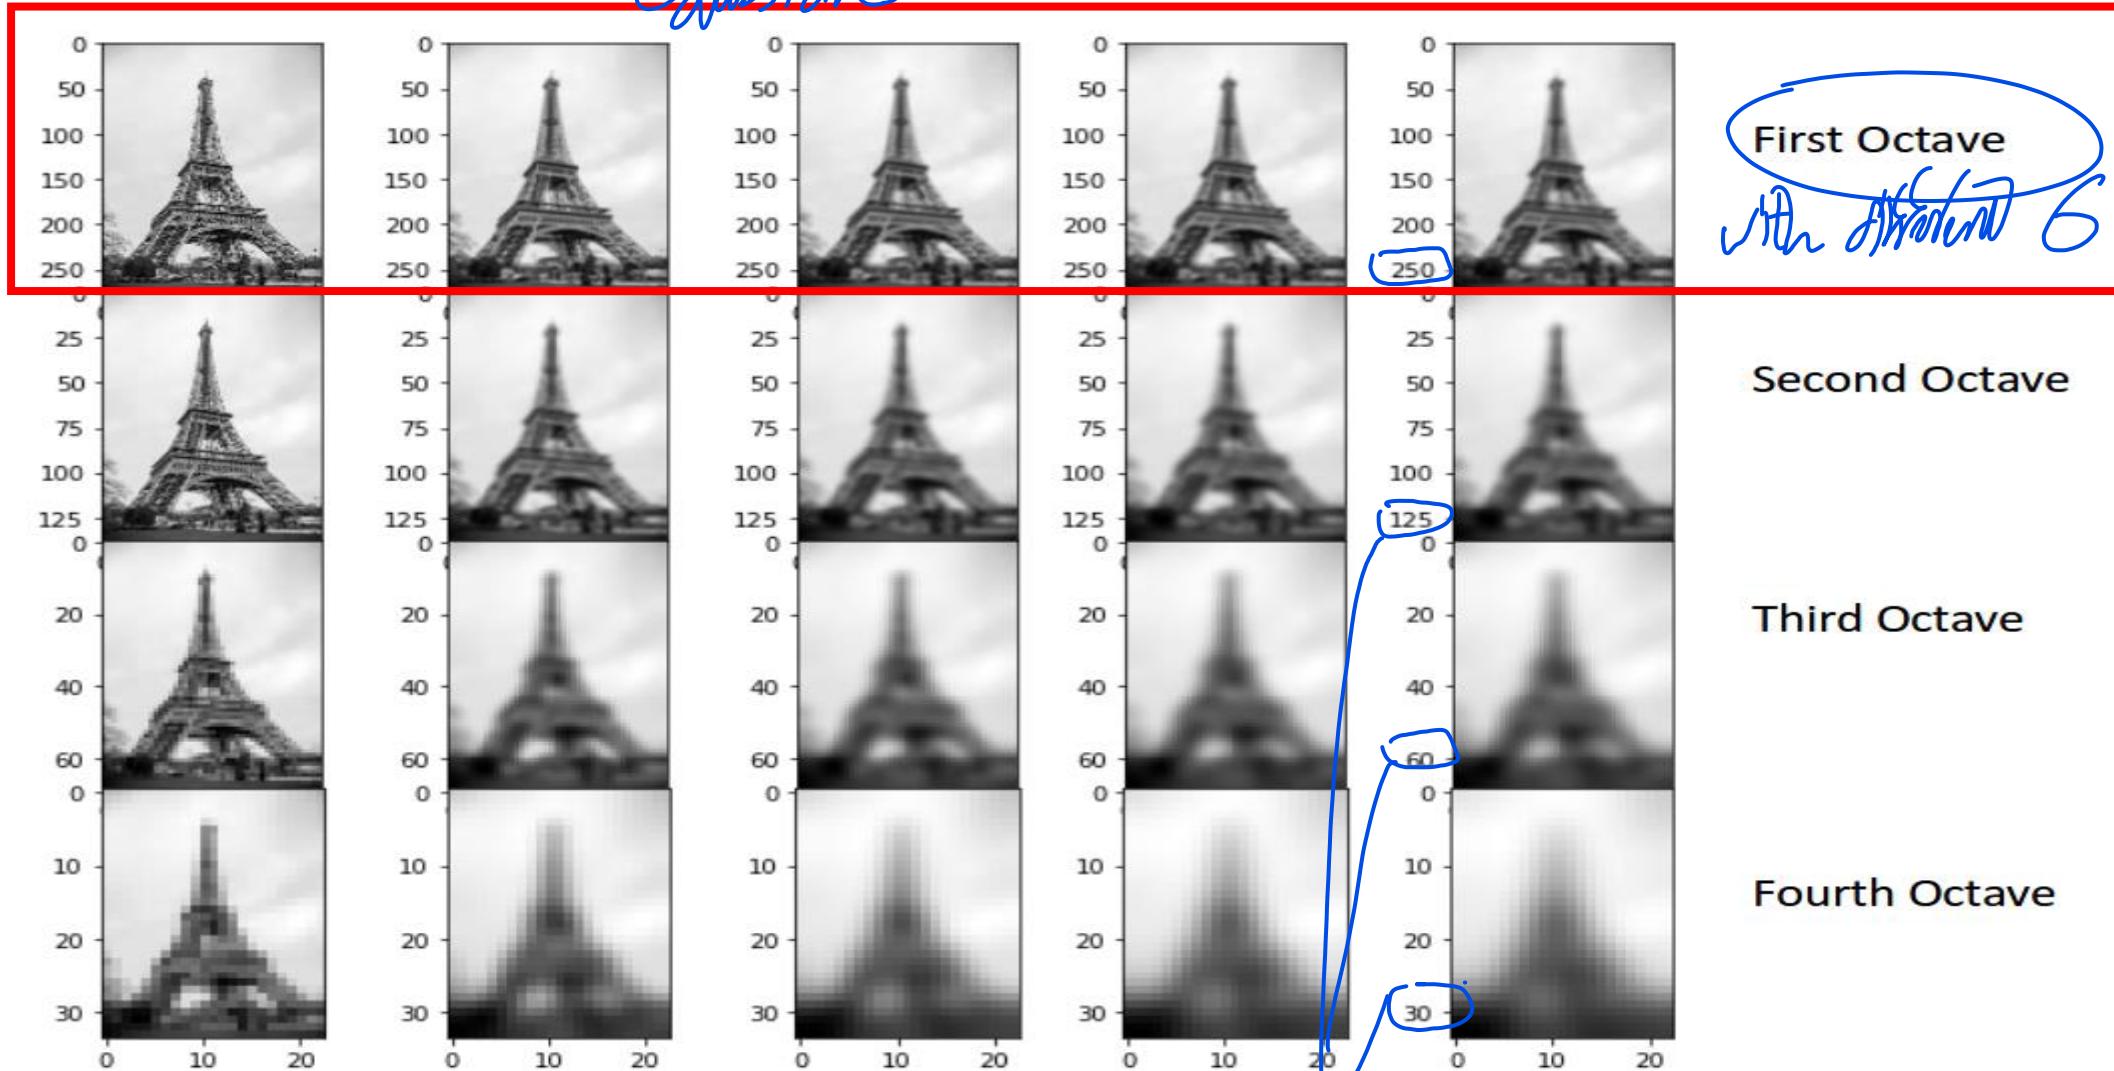


Image courtesy: Lowe40

The ideal number of octaves **should be 4**, and for each octave, the number of blur images should be 5.

4 octaves contains 5 blur images

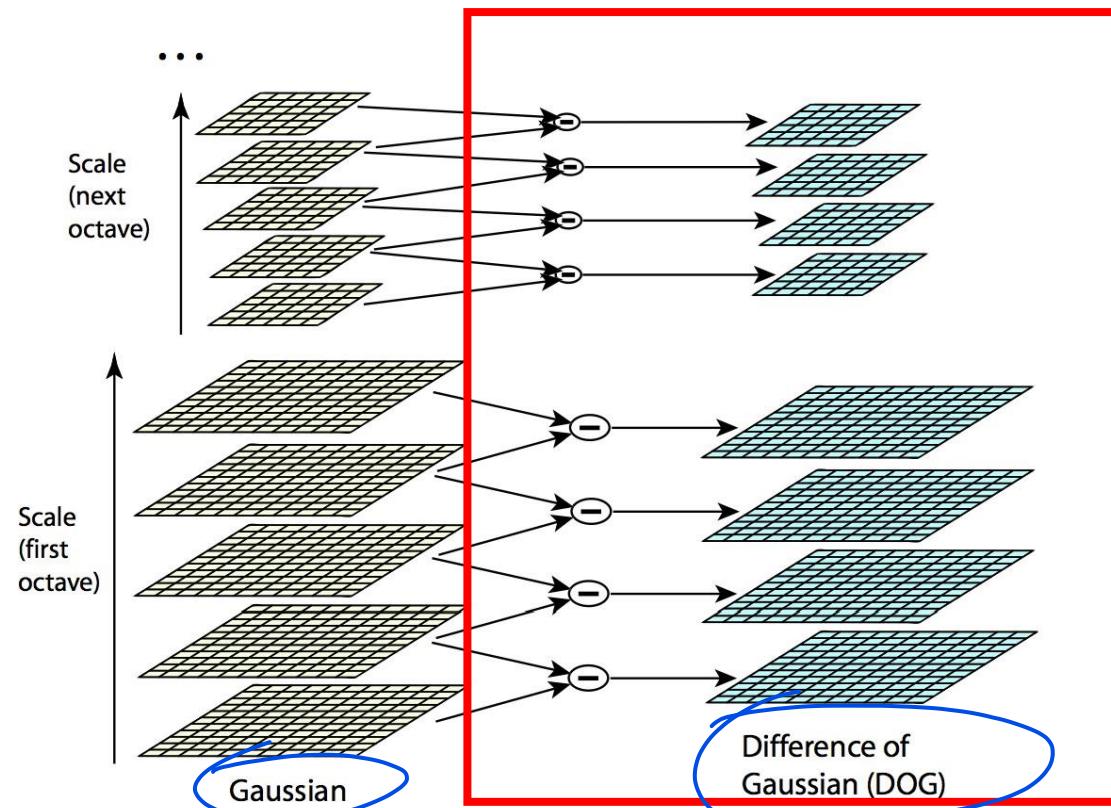
Gaussians



different 5 scales on each octave

# Step 2: Difference of Gaussians

- DoG is a feature enhancement algorithm that creates another set of images, for each octave, by subtracting every image from the previous image on the same scale.



Step 2: Difference of Gaussians  
Subtracting one image from another  
Difference between neighboring scales

## Step 2:

- § Blurring filters out high-frequencies (noise)
- § Subtracting differently blurred images from each other only keeps the frequencies that lie between the blur level of both images
- § DoG acts as a **band-pass** filter

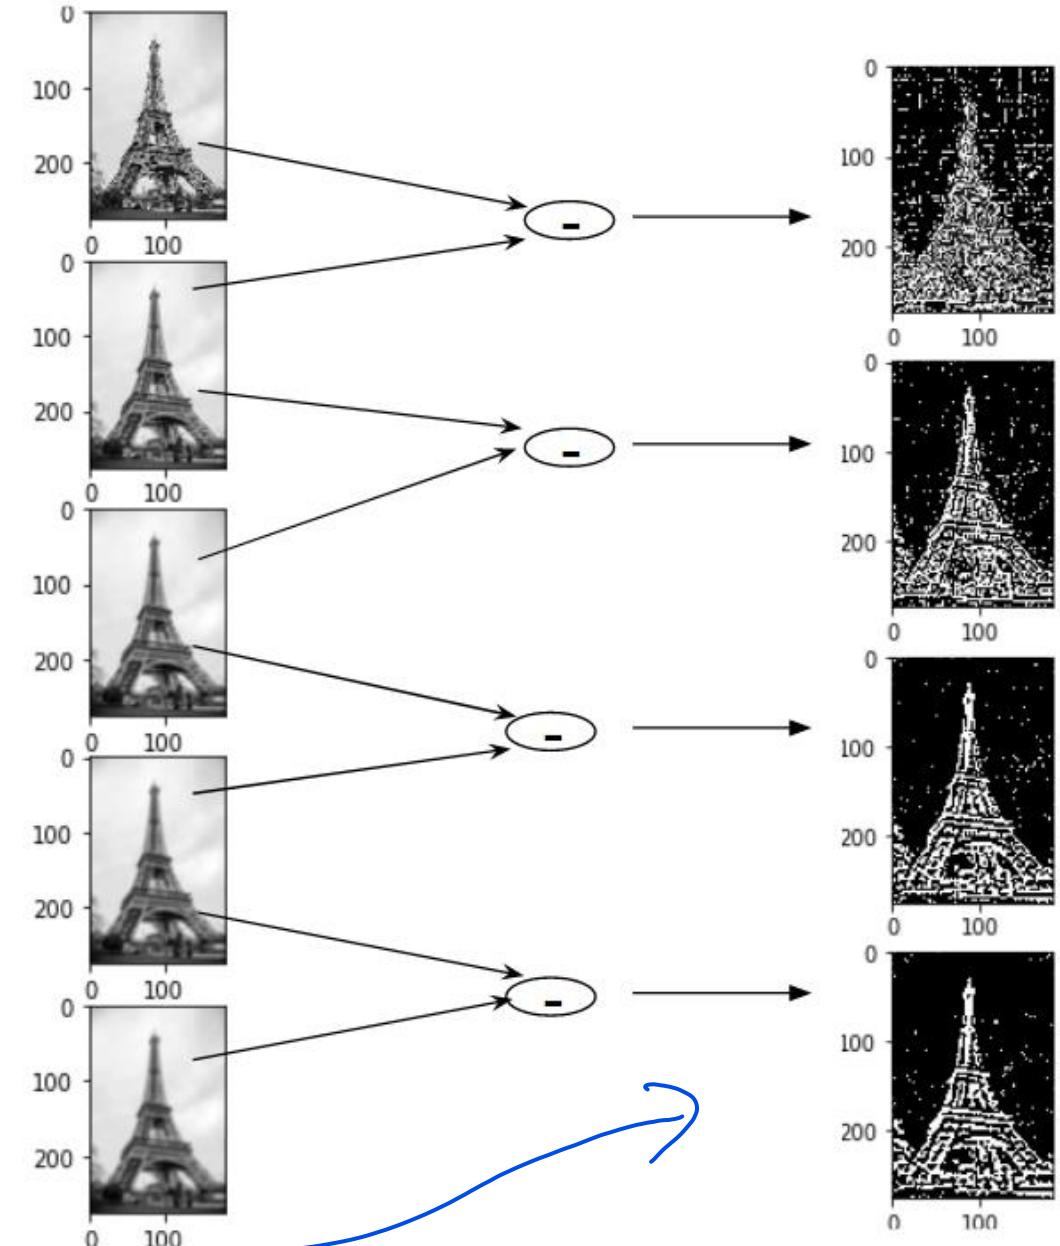
new, girl's sweater has been blurred

## Step 2: Illustration

- On the left, we have 5 images, all from the first octave (thus having the same scale). Each subsequent image is created by applying the Gaussian blur over the previous image.
- On the right, we have four images generated by subtracting the consecutive Gaussians.

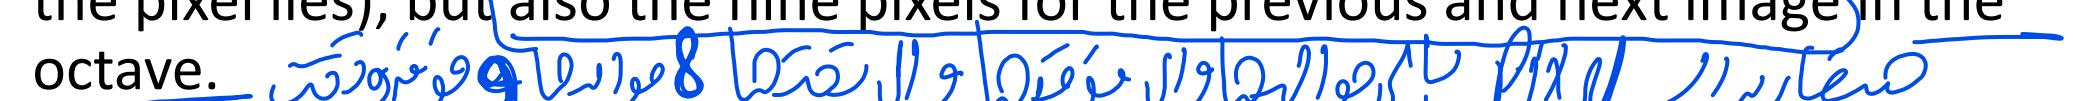
↙ Step ↘ always

→ Increases visibility of corners, edges, and other detail present in the image



## Step 2: Find Extrema

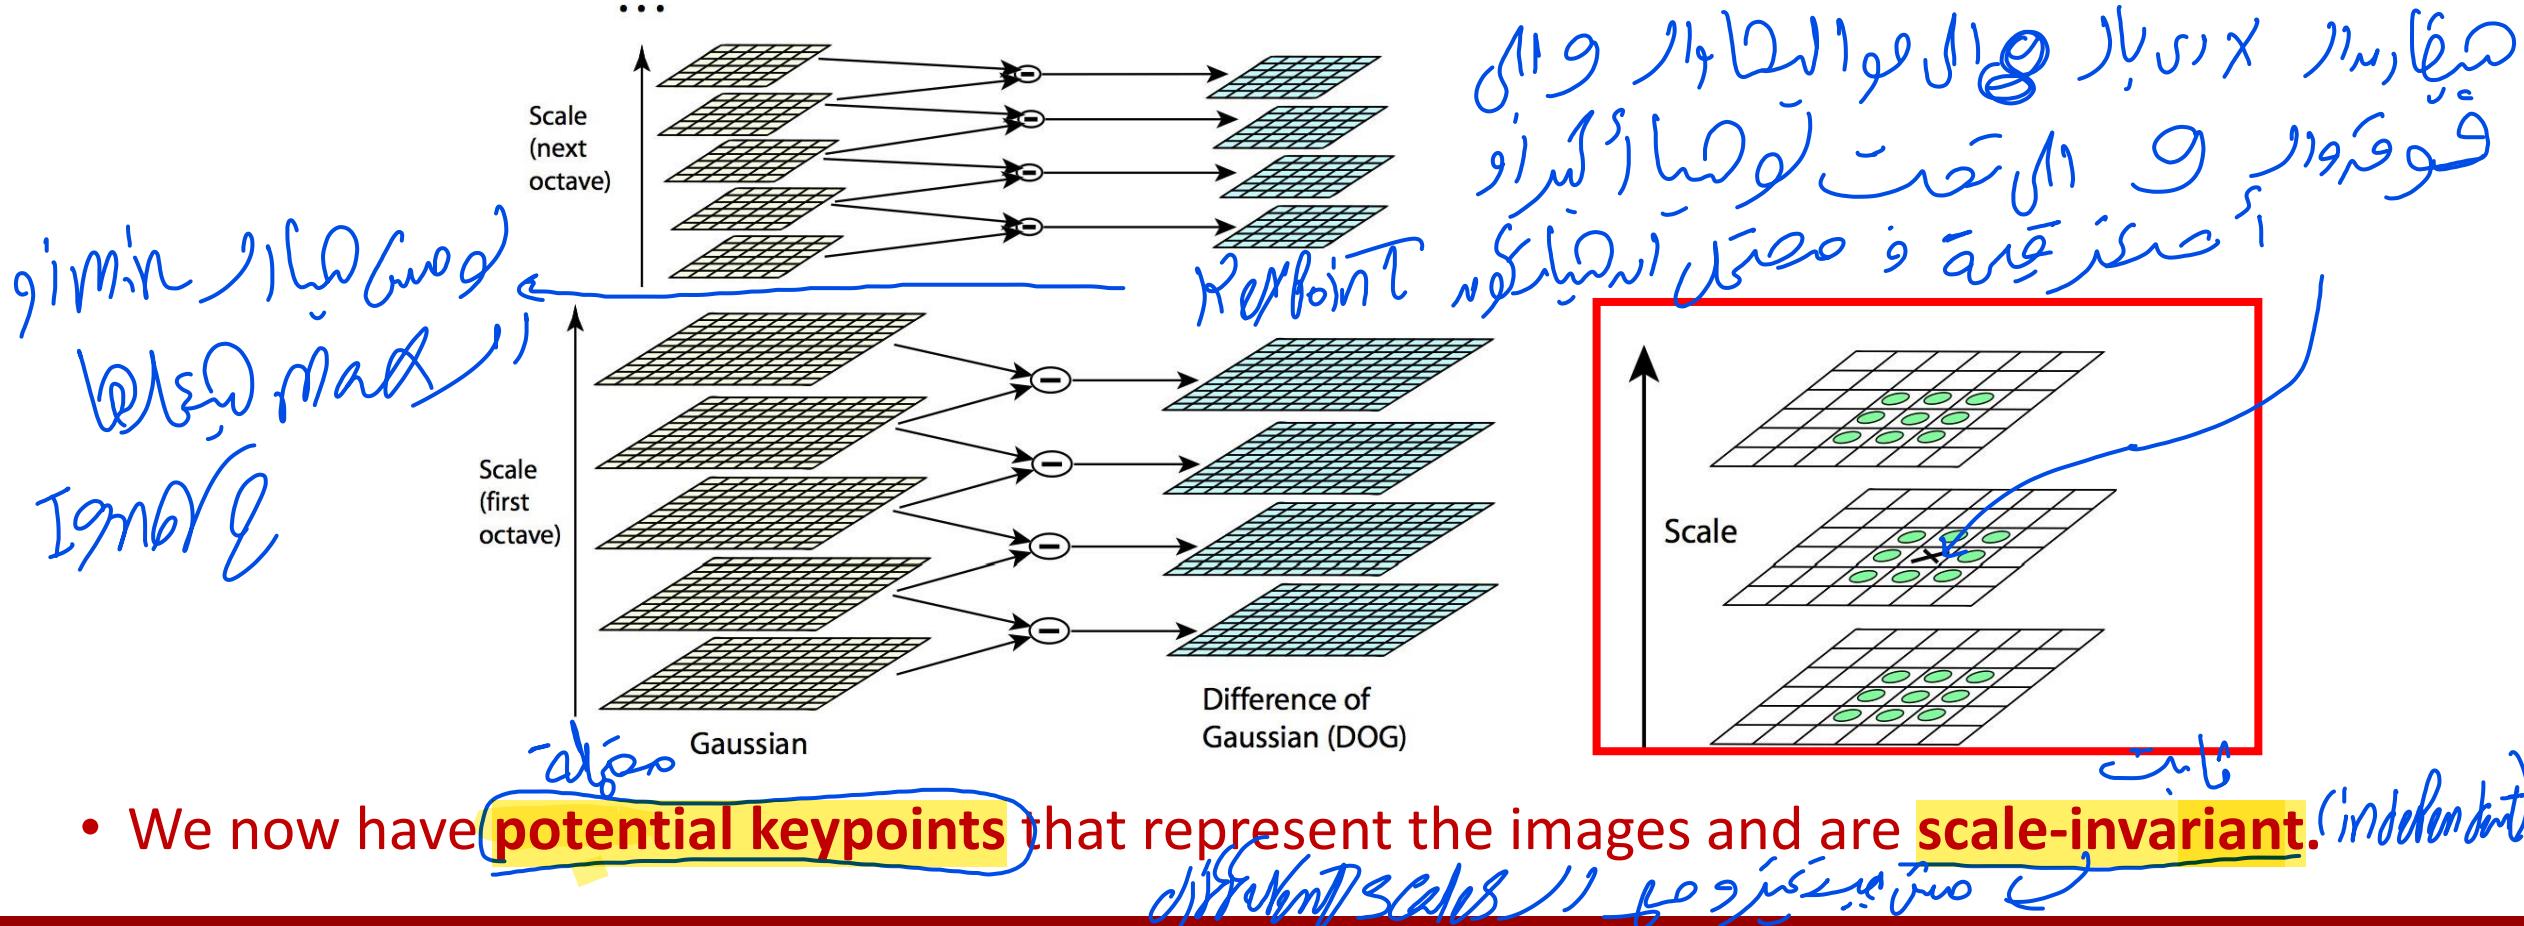
(min or max) Disgusting point of view  
y axis is pointing down left few

- This step is to find the important keypoints from the image that can be used for feature matching.
- The idea is to find **extrema**, i.e., the local maxima and minima for the images.
- To locate the local maxima and minima (extrema):
  - we go through every pixel in the image and compare it with its neighboring pixels.
  - Neighboring not only includes the surrounding pixels of that image (in which the pixel lies), but also the nine pixels for the previous and next image in the octave.  


This means that every pixel value is compared with 26 other pixel values to find whether it is the local maxima/minima.

## Step 2: Find Extrema

- For example, in the below diagram, we have three images from the first octave. The pixel marked  $x$  is compared with the neighboring pixels (in green) and is selected as a keypoint if it is the highest or lowest among the neighbors.



## Step 3: Extrema suppression (Keypoint Localization)

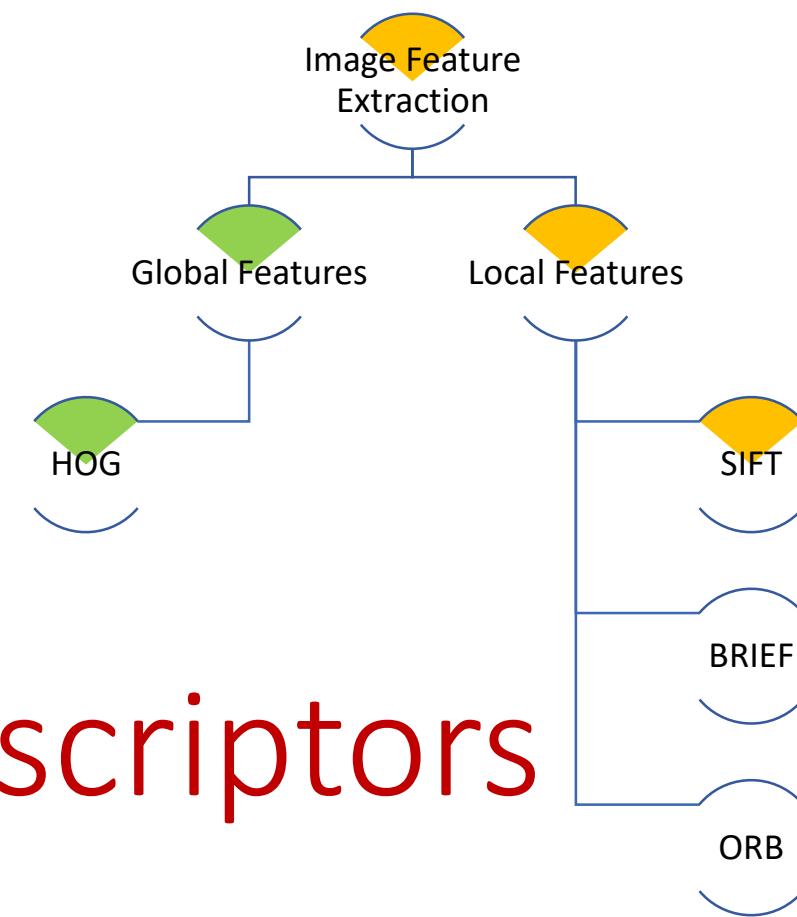
- § The DoG finds blob-like and corner-like image structures but also leads to strong responses along edges
- § **Corners and blobs are good keypoints, but edges are bad for matching.**

→ Hence, we will eliminate the keypoints that may not be robust to noise such as low contrast, and edges.

low contrast), edges (non-maxima  
(PT with links))

## Step 3:

- To deal with the low contrast keypoints, a second-order Taylor expansion is computed for each keypoint. If the resulting value is less than 0.03 (in magnitude), we reject the keypoint.
  - To deal with edge keypoints, a second-order Hessian matrix is used to identify such keypoints.
- we have accurate keypoints to represent the image features with scale-invariant.

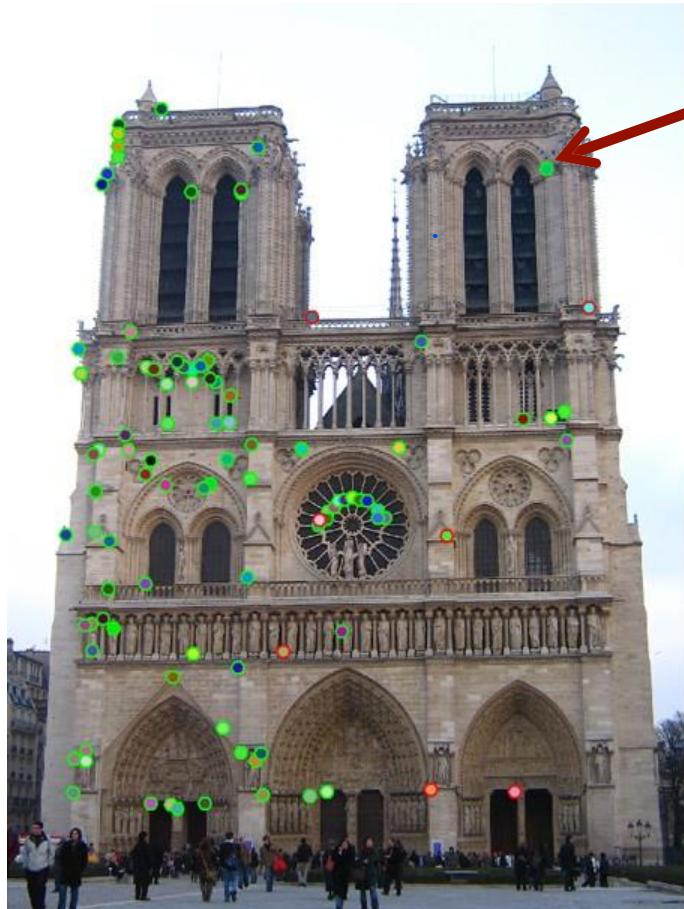


# Local features: Keypoints & Descriptors

SIFT Descriptor

Descritor  
Keypoint

# Keypoint Done. What about a Descriptor?



**keypoint**  
**descriptor** at  
the keypoint

$$f = \begin{bmatrix} 0.02 \\ 0.04 \\ 0.1 \\ 0.03 \\ 0 \\ \dots \end{bmatrix}$$

# Local Descriptors

## § Features: Describing a keypoint

- § **SIFT** – Scale Invariant Feature Transform
- § **SURF** – Speeded-Up Robust Features
- § **BRIEF** – Binary Robust Independent Elementary Features
- § **ORB** – Oriented FAST Rotated BRIEF
- § **BRISK** – Binary Robust Invariant Scalable Keypoints
- § **FREAK** – Fast REtinA Keypoint

وكلها ملخص.

local descriptors

only 0,1 values Binary Descriptors

# SIFT Considers the Distribution of Gradients Around Keypoints

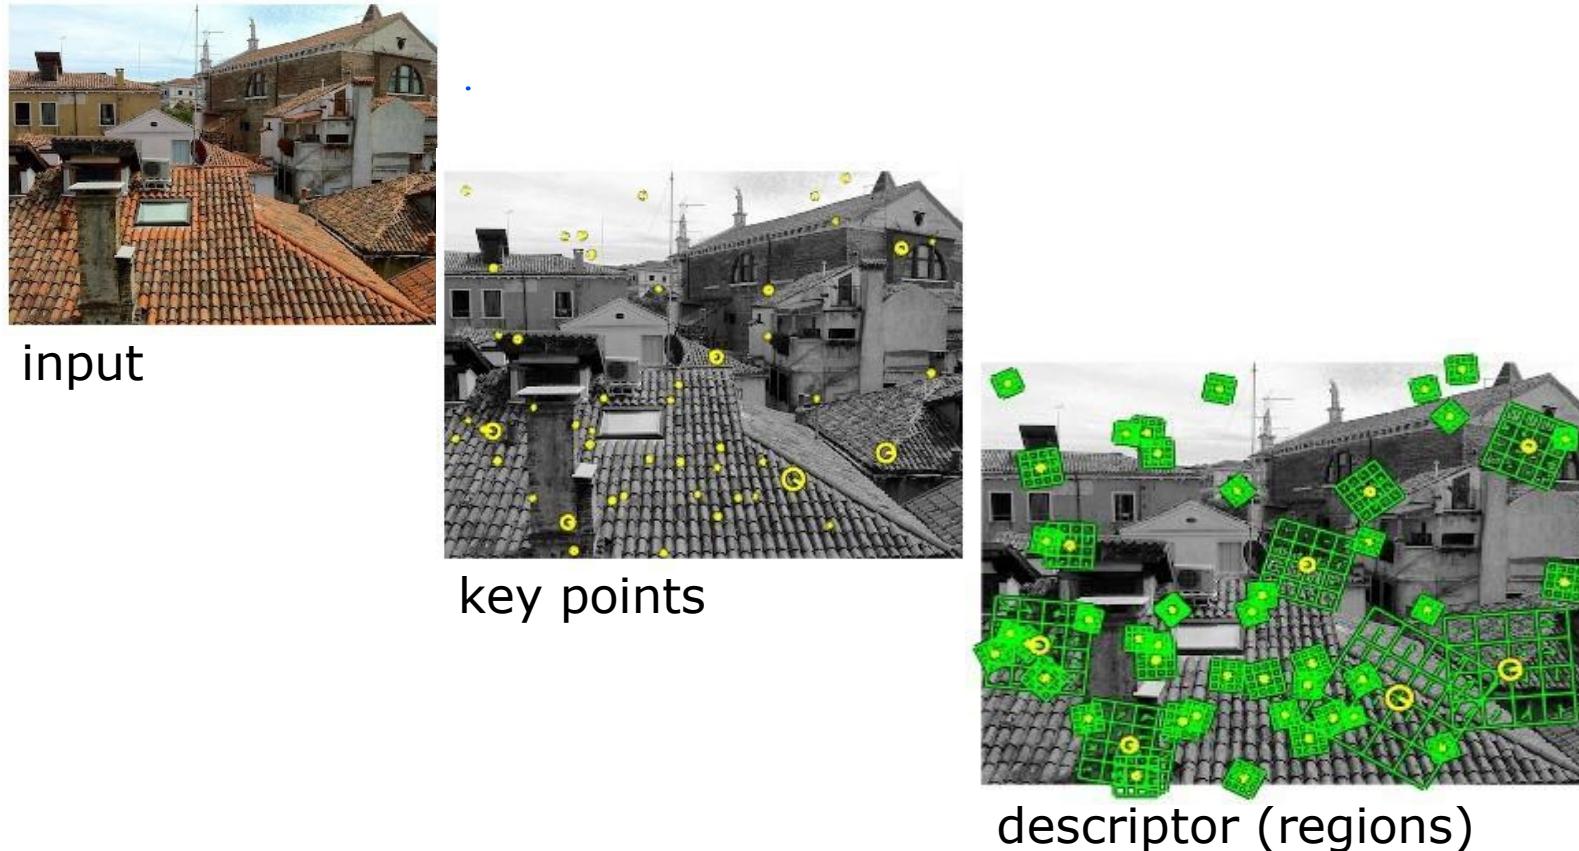


Image courtesy: Vedaldi and Fulkerson 19

# SIFT Descriptor

§ Image content is transformed into features that are **invariant to**

- § image translation,
- § rotation, and
- § scale

*independent*

*as columns are 'that's independent'*  
*different scales*

§ They are **partially invariant to**

- § illumination changes and
- § affine transformations and 3D projections

*so it's not going to change*  
*affine*

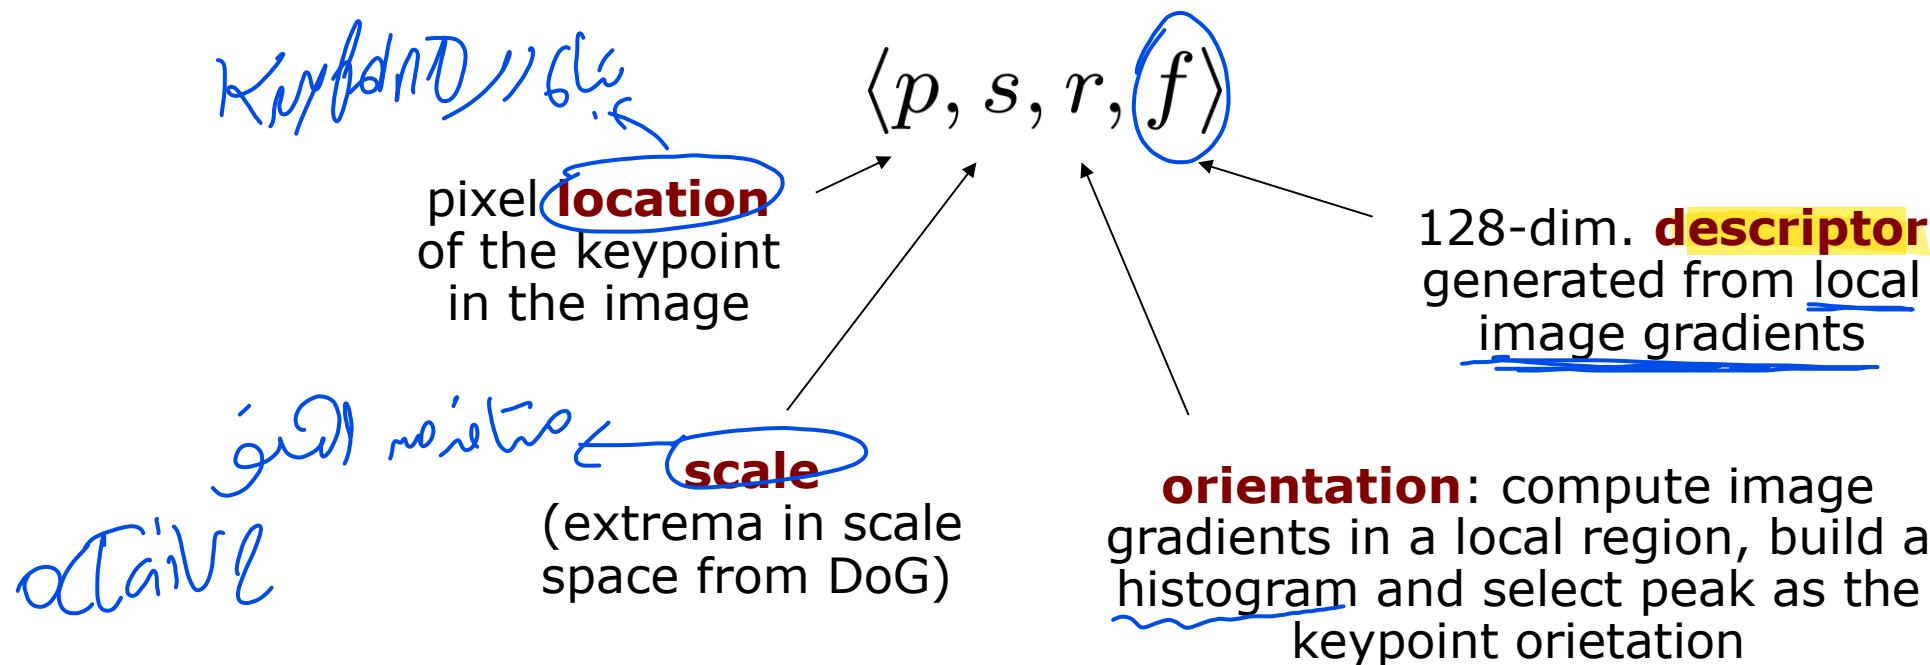
§ Suitable for **detecting visual landmarks**

- § from different angles and distances
- § with a different illumination

*both*

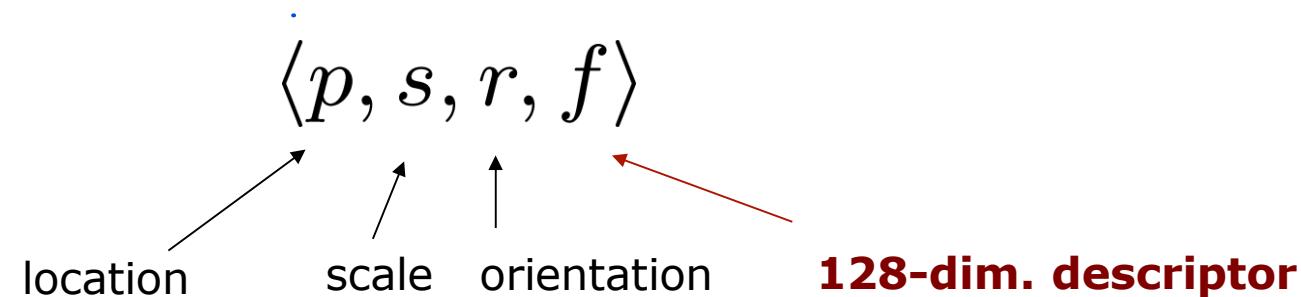
# SIFT Features

- A SIFT feature is given by a vector computed at a local extreme point in the scale space



# SIFT Features

- A SIFT feature is given by a vector computed at a local extreme point in the scale space



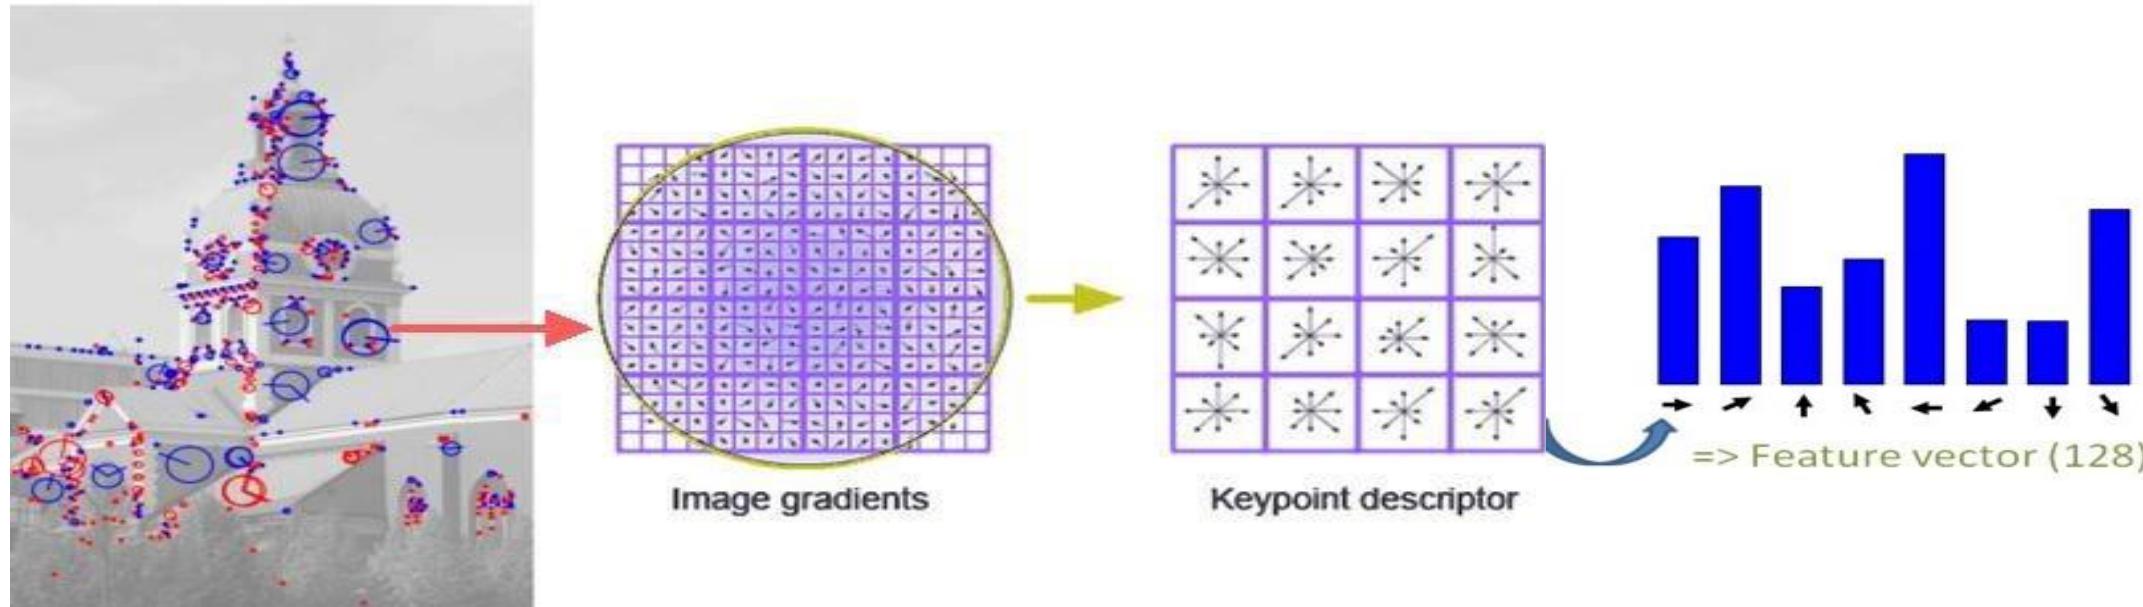
view-point dependent

mainly independent

# SIFT Descriptor Steps

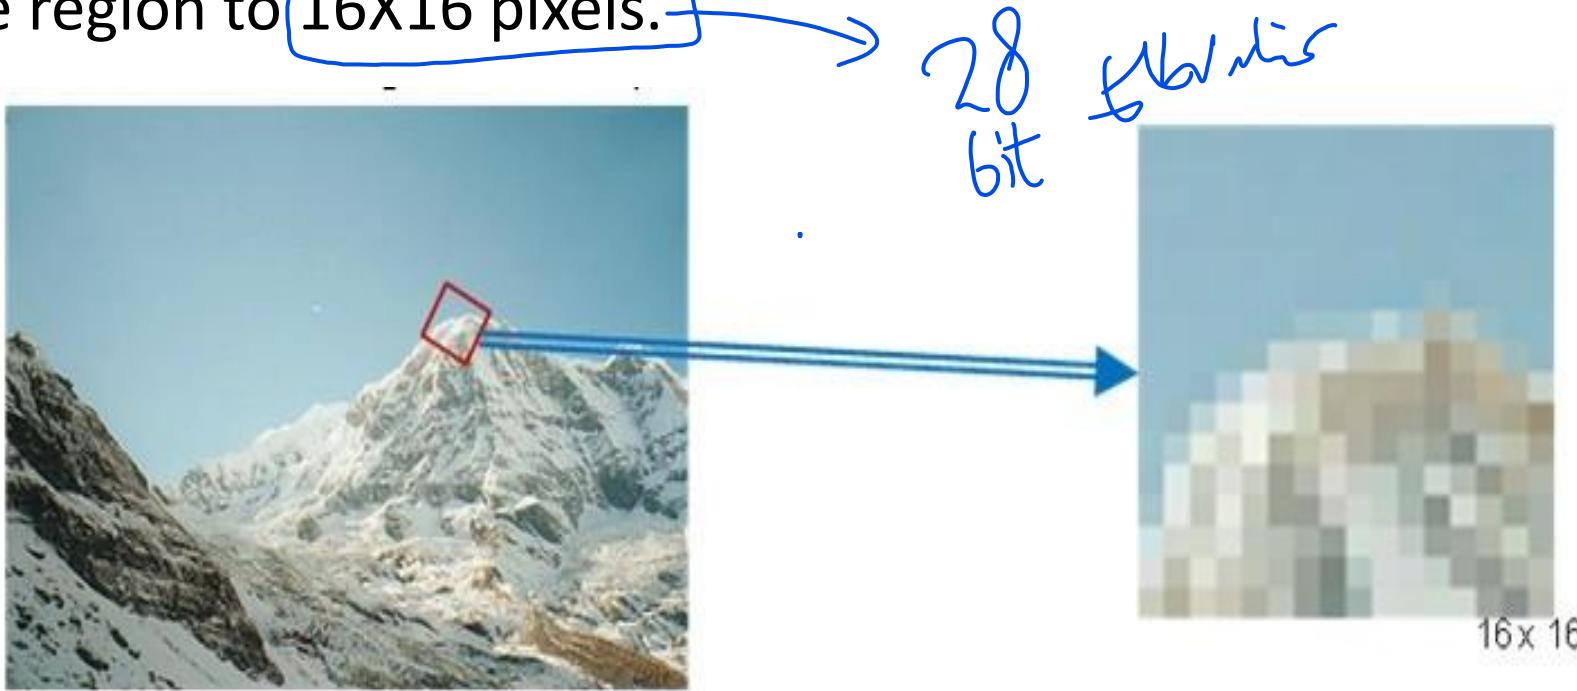
- Step1: Compute image gradients in local  $16 \times 16$  area at the selected scale
- Step2: Create an array of orientation histograms
- Step3:  $8$  orientations  $\times$   $4 \times 4$  histogram array =  $128$  dimensions  
(yields best results)

$$8 \times 4 \times 4 = 128 \times 1 \leftarrow \text{for 1 pixel}$$



# Step 1: Compute image gradients in local $16 \times 16$ area at the selected scale

- For a given keypoint, warp the region around it to orientation and scale and resize the region to  $16 \times 16$  pixels.

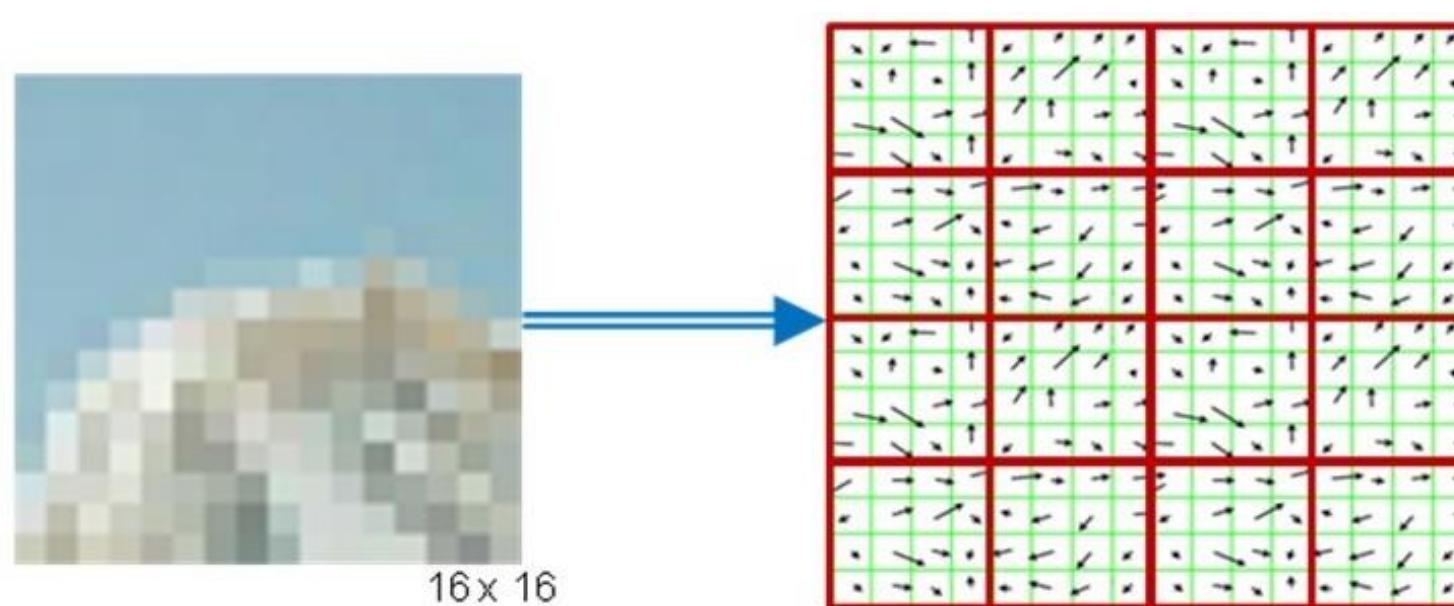


- Since we use the surrounding pixels, the descriptors will be partially invariant to illumination or brightness of the images.

## Step 1:

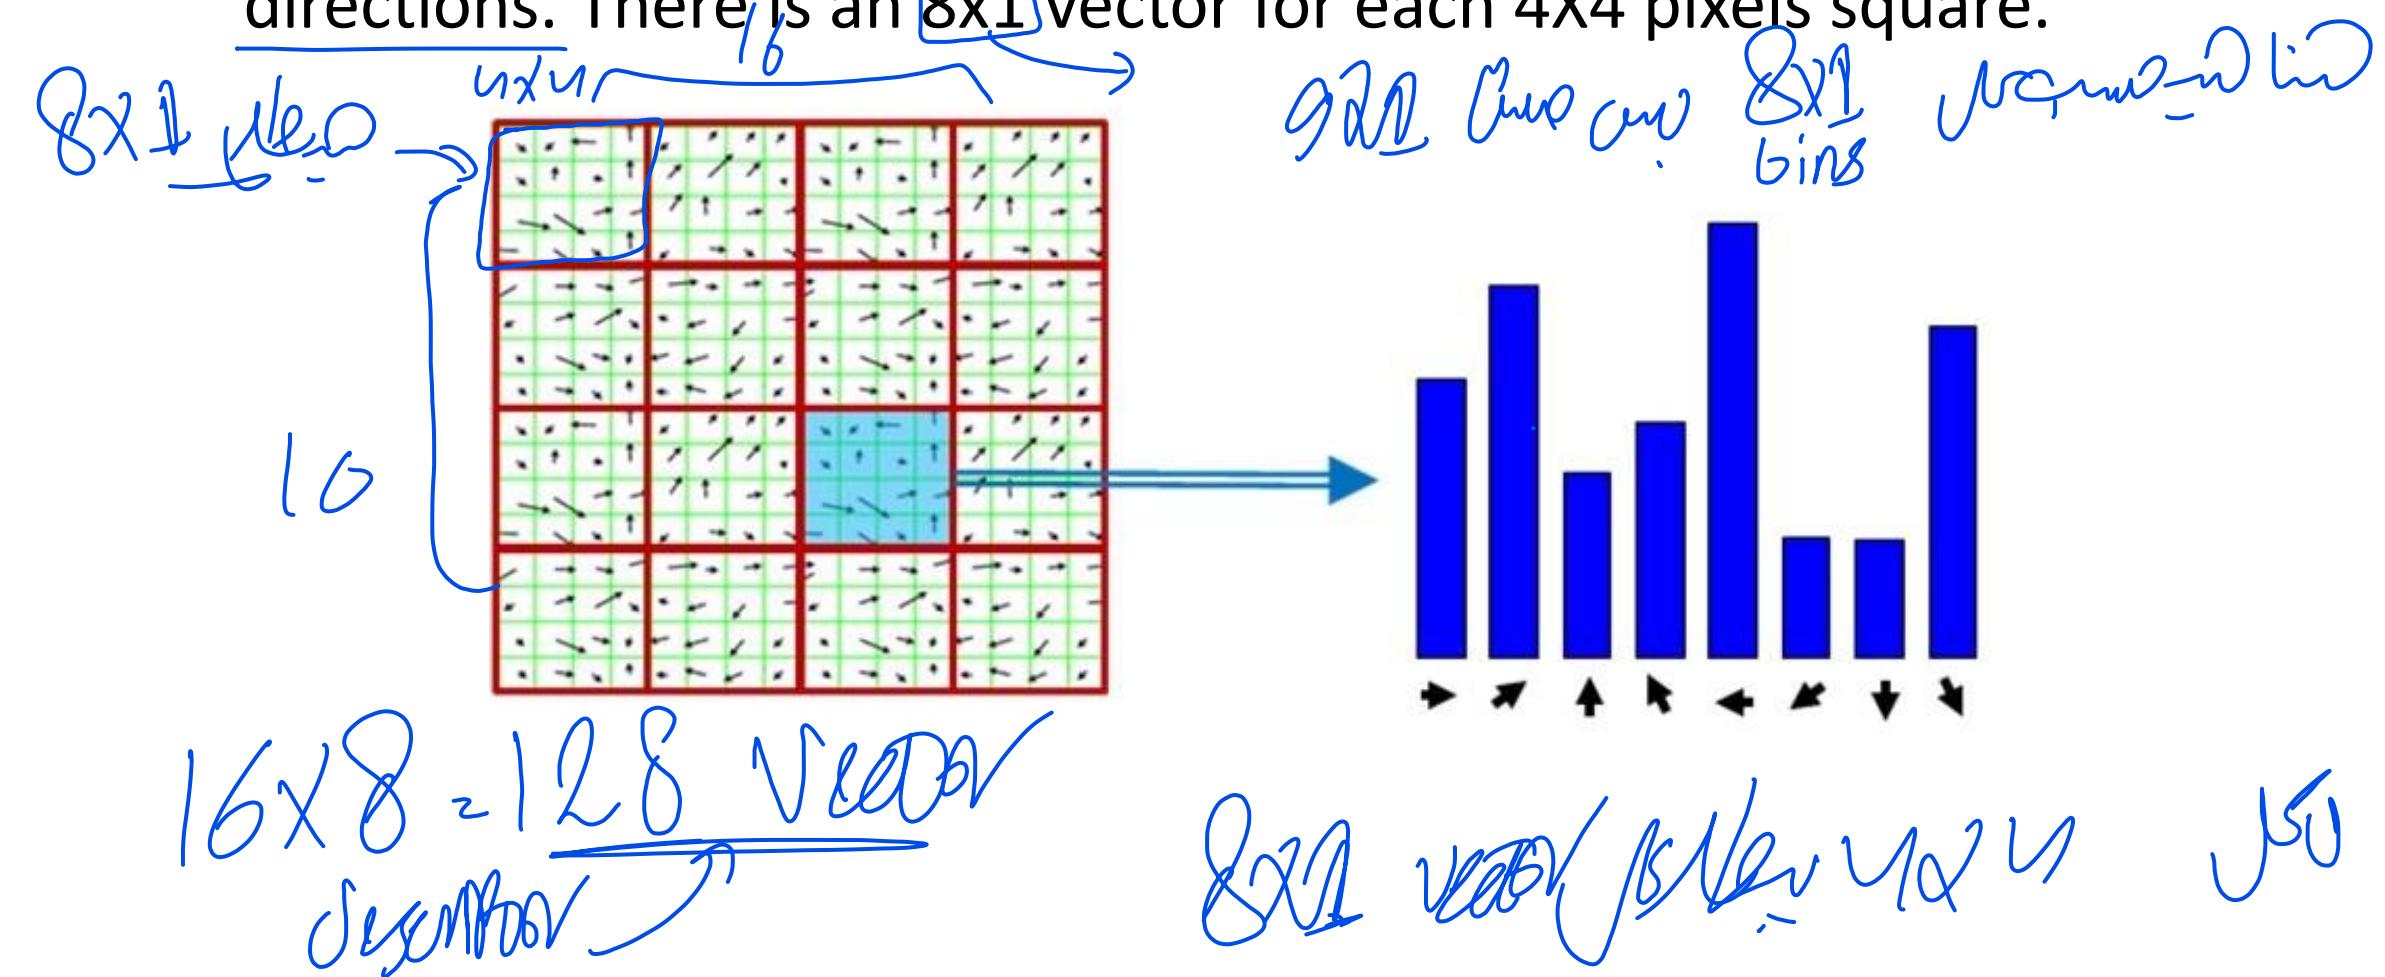
HOG feature

- Compute the gradients for each pixel (orientation and magnitude) to generate a unique fingerprint for this keypoint.
- Then, divide the pixels into 16, 4X4 pixels squares.



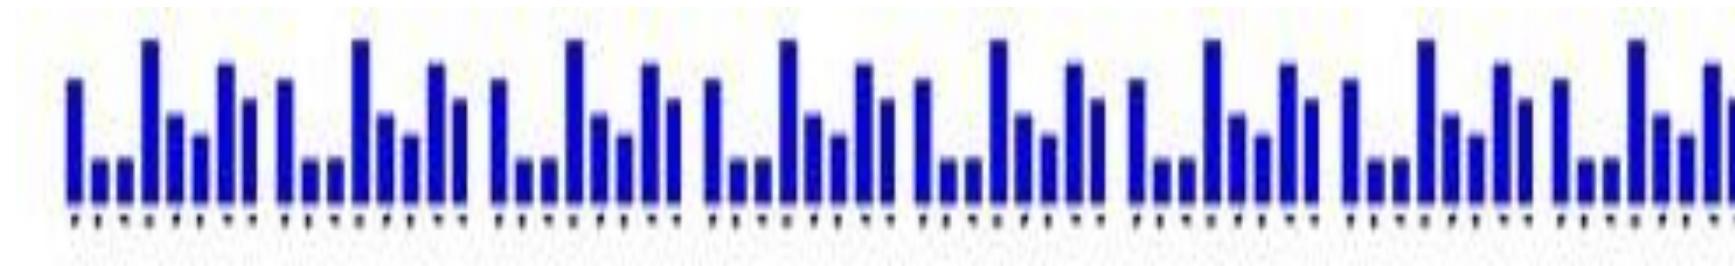
## Step 2: Create an array of orientation histograms

- For each square, compute gradient direction histogram over 8 directions. There is an 8x1 vector for each 4X4 pixels square.

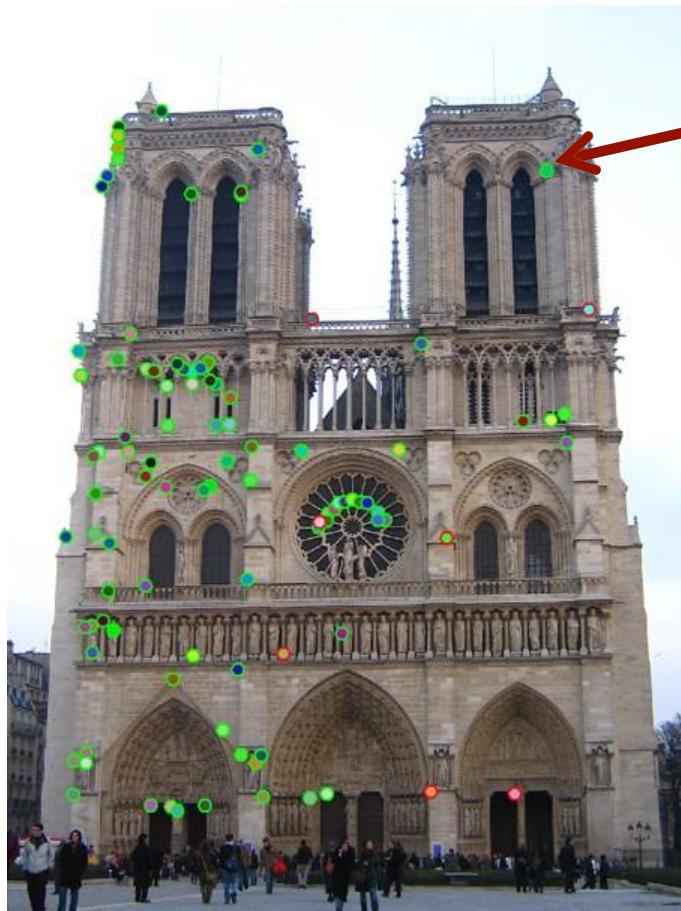


## Step 3: 8 orientations x 4x4 histogram array

- Concatenate the histograms to obtain a 128 (16\*8) dimensional feature vector:



# SIFT Approach Done!



**keypoint**  
(via DoG)

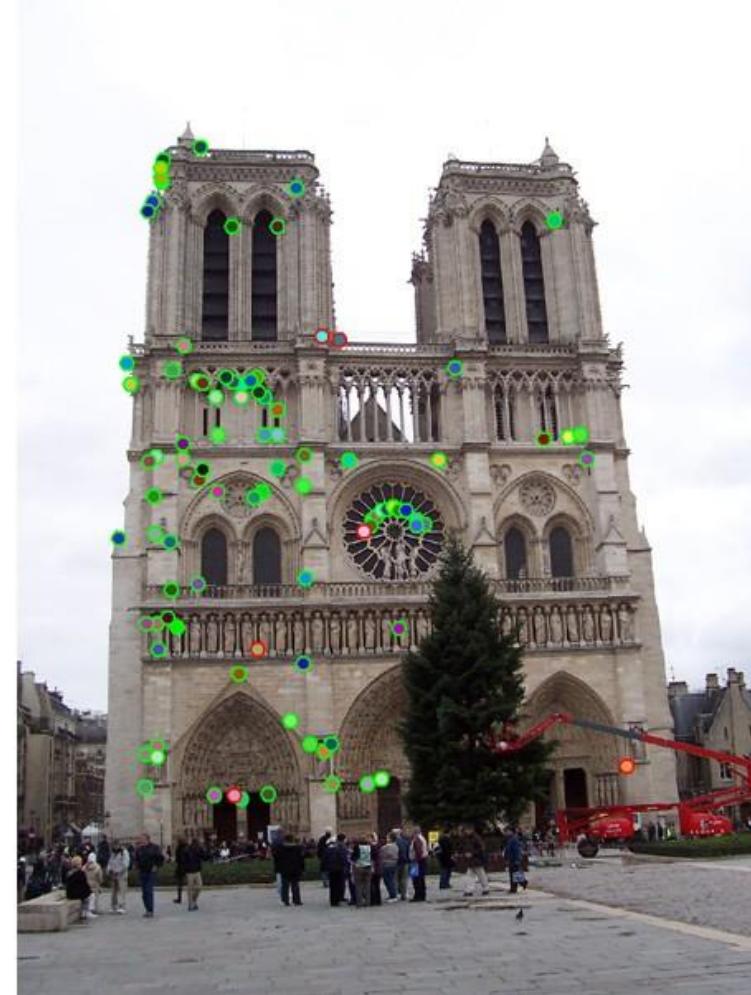
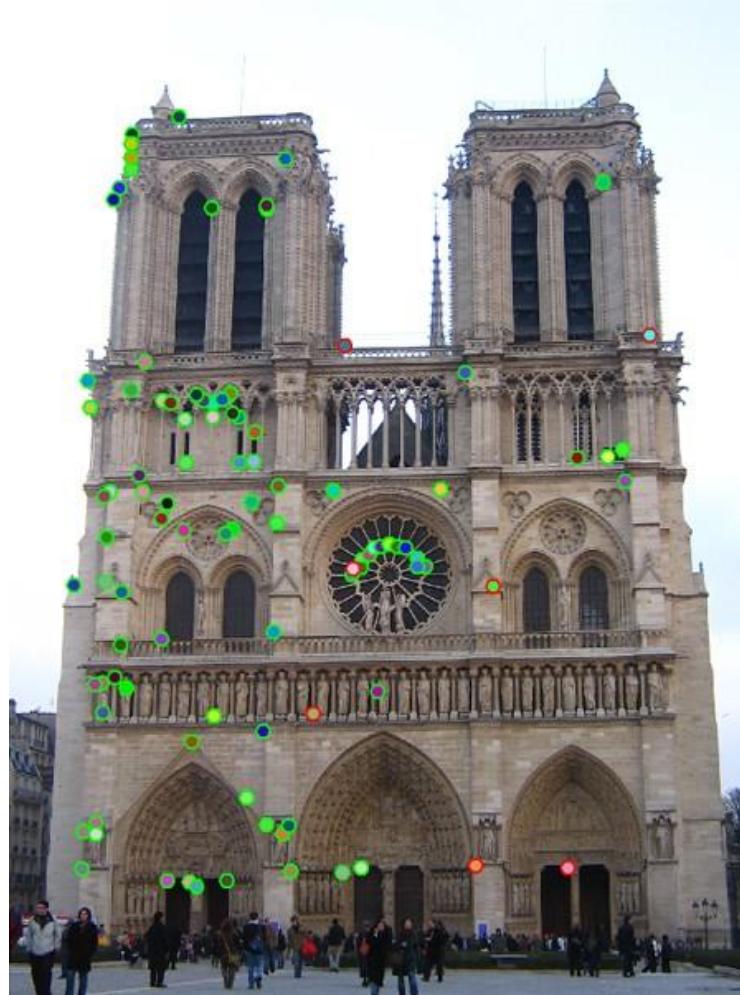
**descriptor** (via  
gradient histogram)

$$f = \begin{bmatrix} 0.02 \\ 0.04 \\ 0.1 \\ 0.03 \\ 0 \\ \dots \end{bmatrix}$$

# How To Match Them?

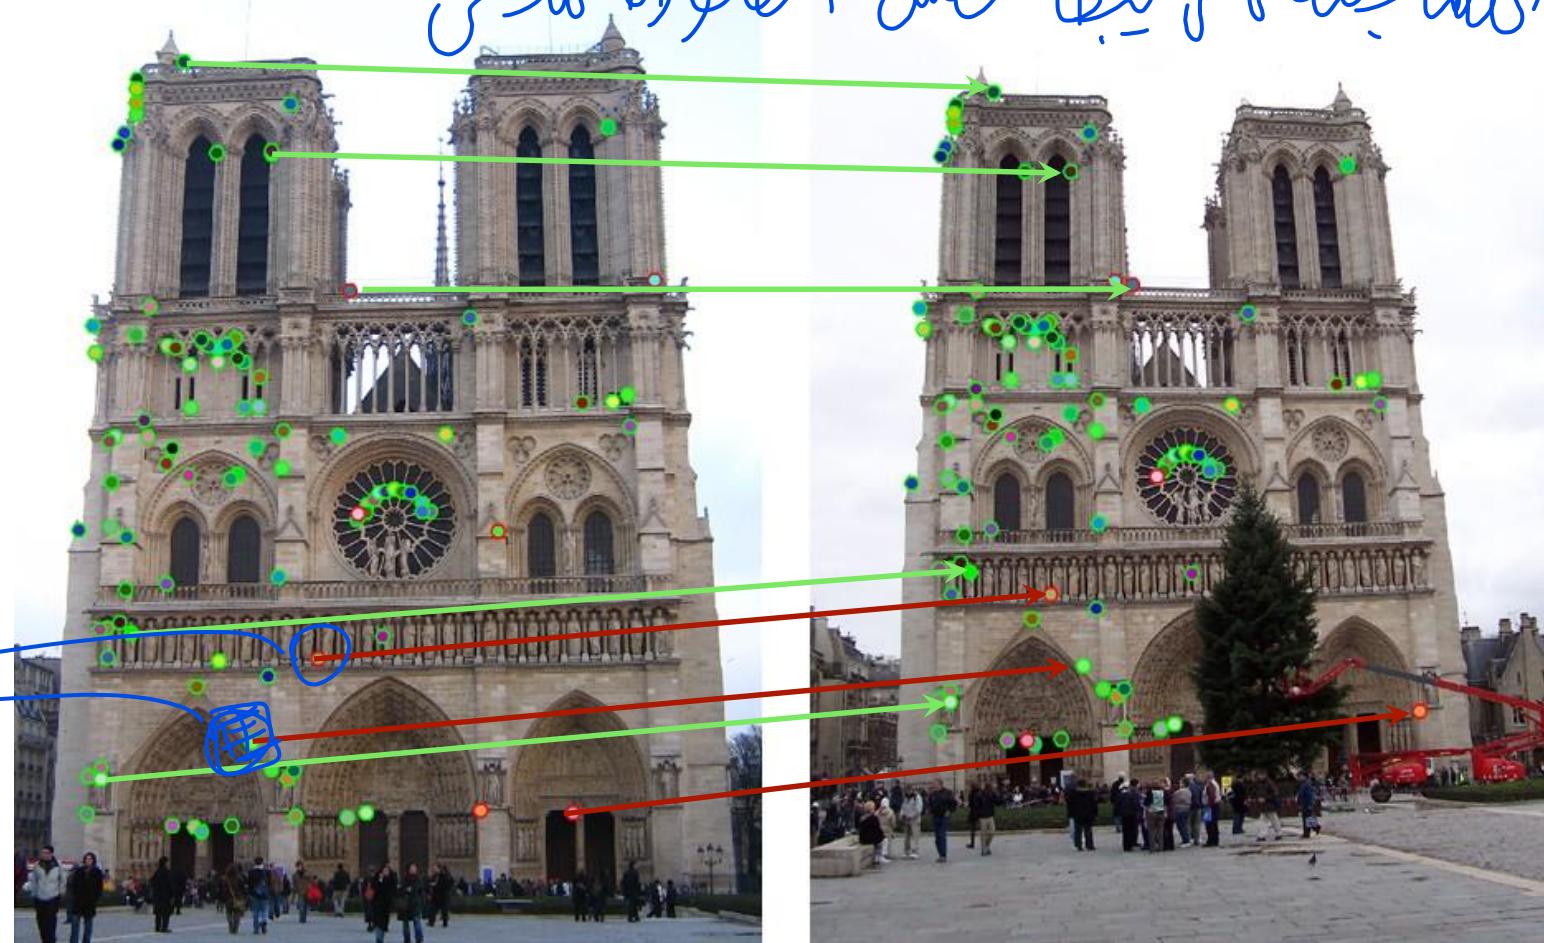
it's Sal

Neural



# Based on Descriptor Difference

Match سیستم مبتنی بر تفاوت دسکرپتور است که از تفاوت دسکرپتورهای متناظر در دو تصویر برای پیدا کردن مطابقت می‌شود.



# Lowe's Ratio Test

سلویس

§ 3 Steps test to eliminate ambiguous matches for a query

feature  $q$ : *أقرب المصفوفات في المصفوفة المعرفة*

1. Step: Find the closest two descriptors to  $q$ , called  $p_1$  and  $p_2$  based on the Euclidean distance  $d$

$$d(p, q) = \sqrt{\sum_{i=1}^n (q_i - p_i)^2}$$

2. Step: Test if the distance to the best match is smaller than a threshold:

$$d(q, p_1) < T$$

3. Step: Accept match only if the best match is substantially better than second:

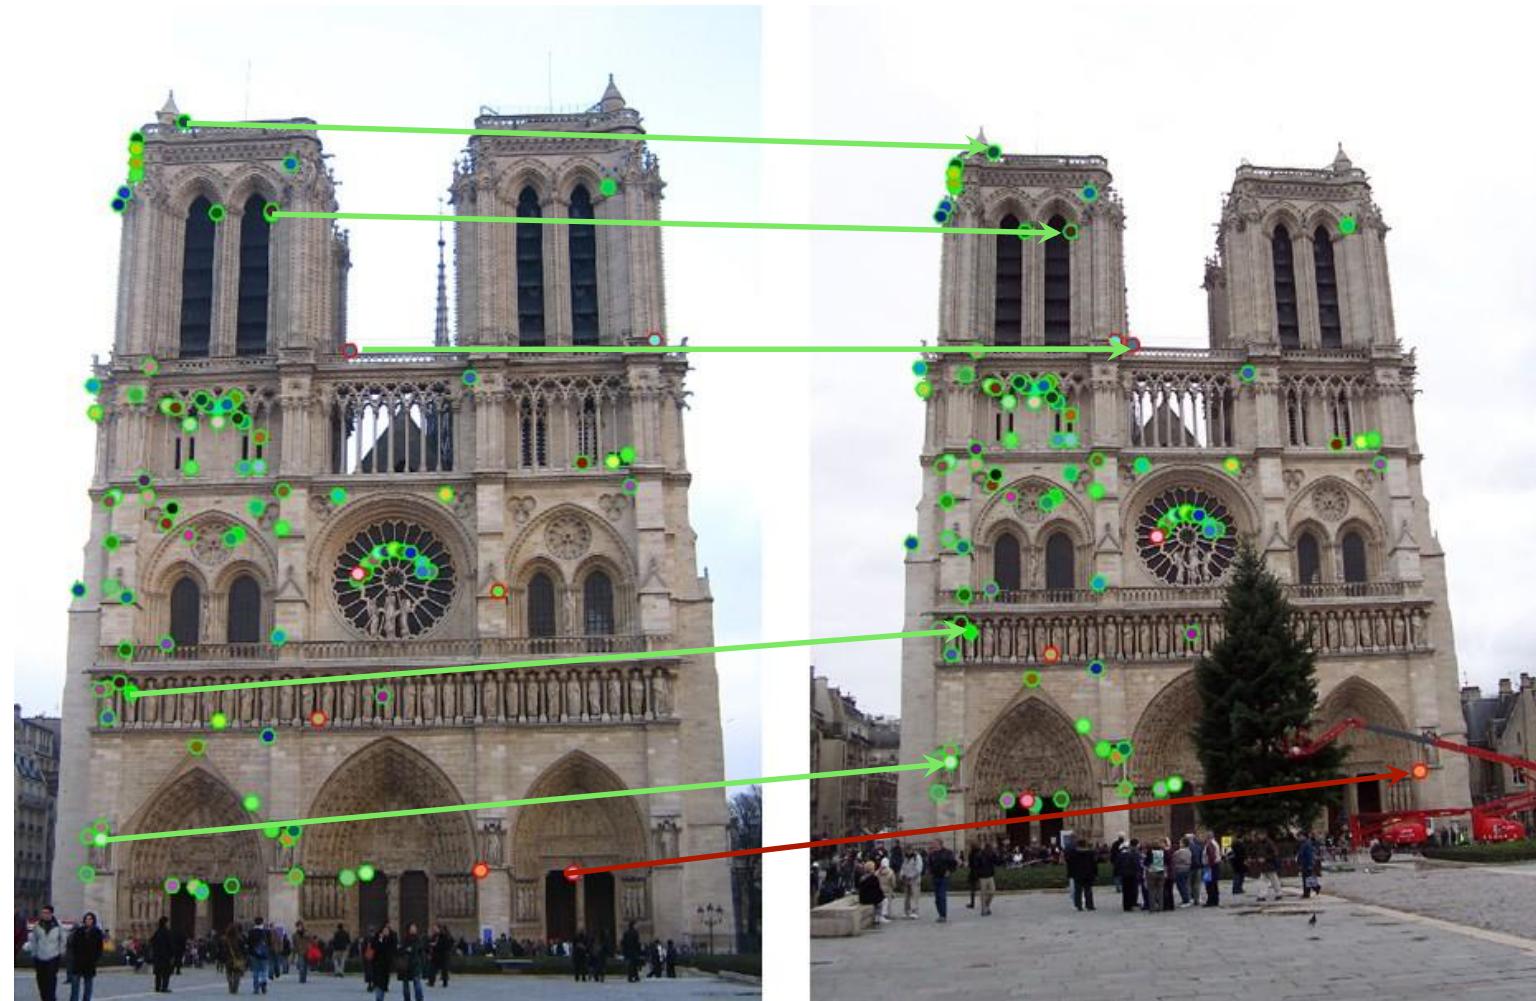
$$\frac{d(q, p_1)}{d(q, p_2)} < \frac{1}{2}$$

Ignore? *ما هي الخطأ؟*

Ignore? *ما هي الخطأ؟*

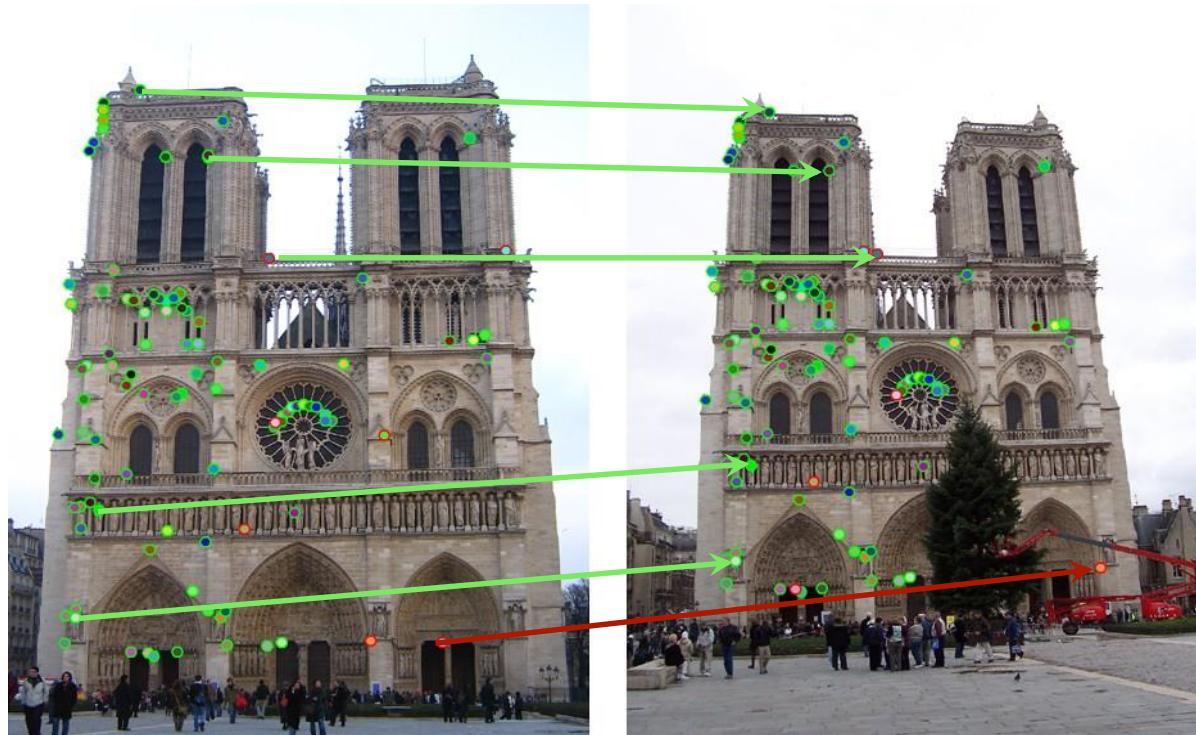
# Based on RatioTest

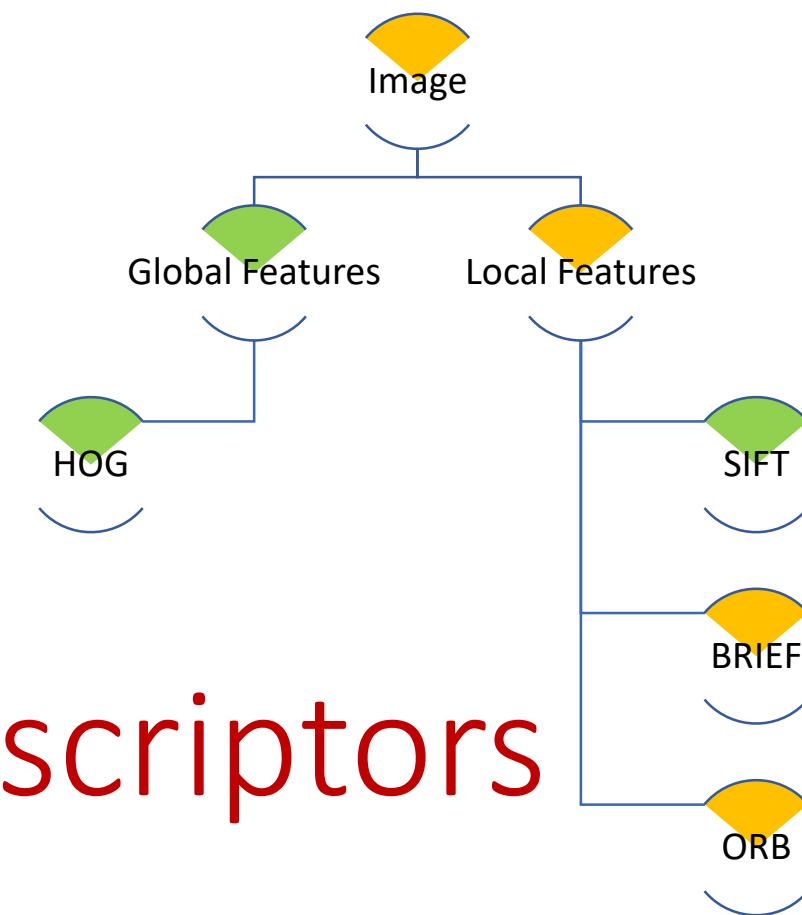
62 دیکی نیکی پیشنهاد



# Outliers

- § Lowe's Ratio test works well
- § There will still remain few outliers
- § Outliers require an extra treatment





# Local features: Keypoints & Descriptors

**Binary Descriptors: BRIEF, ORB**

# Why Binary Descriptors? *plz, i, g, l, w, 2, i*

- § **Complex features** such as SIFT work well and are a gold standard
- § SIFT is expensive to compute
- § Binary descriptors aim at generating **small binary strings** that are easy to compute and compare

# Key Advantages of Binary Descriptors

- § **Compact descriptor**
  - The number of pairs gives the length in bits
- § **Fast to compute**
  - Simply intensity value comparisons
- § **Trivial and fast to compare**
  - Hamming distance

$$d_{\text{Hamming}}(B_1, B_2) = \text{sum}(\text{xor}(B_1, B_2))$$

# Examples of Binary descriptors

- **BRIEF** – Binary Robust Independent Elementary Features
- **ORB** – Oriented FAST Rotated BRIEF
- **BRISK** – Binary Robust Invariant Scalable Keypoints
- **FREAK** – Fast REtinA Keypoint

# Summary

- § Keypoints and descriptor together define common visual features
- § Keypoint defines the location
- § Descriptor describes the appearance الляр (ل)
- § Several descriptors operating on gradient histograms (SIFT, SURF, ...)
- § Binary descriptors for efficiency (BRIEF, ORB, ...)

# DataCamp Courses

## CAREER TRACK: Machine Learning Scientist with Python

**Phase 1: (5 points) due Fri. 24 Nov.**

- 1- Supervised Learning with scikit-learn (1)
- 2- Preprocessing for Machine Learning in Python (8)
- 3- Feature Engineering for Machine Learning in Python (10)
- 4- Linear Classifiers in Python (3)

**Phase 2: (5 points) due Fri. 22 Dec.**

- 1- Machine Learning with Tree-Based Models in Python (4)
- 2- Cluster Analysis in Python (6)
- 3- Dimensionality Reduction in Python (7)
- 4- Model Validation in Python (11)

# Thanks